



UNIVERSITY OF COPENHAGEN

# Network Apps: Overview of Socket API

## Network Apps: HTTP & Content Delivery

David Marchant

Based on slides compiled by Marcos Vaz Salles, with adaptiona by Vivek Shah and Michal Kirkedal Thomsen

## Admin

- No lecture or exercises on Wednesday 10<sup>th</sup>
- Will be held on Friday 12<sup>th</sup>
- Same time, same place, different date





## Recap: Key Concepts in Networking



- **Protocols**
  - Speaking the same language
  - Syntax and semantics
- **Layering**
  - Standing on the shoulders of giants
  - A key to managing complexity
- **Resource allocation**
  - Dividing scarce resources among competing parties
  - Memory, link bandwidth, wireless spectrum, paths
- **Naming**
  - What to call computers, services, protocols, ...



# Layer Encapsulation in HTTP



# Layer Encapsulation in HTTP



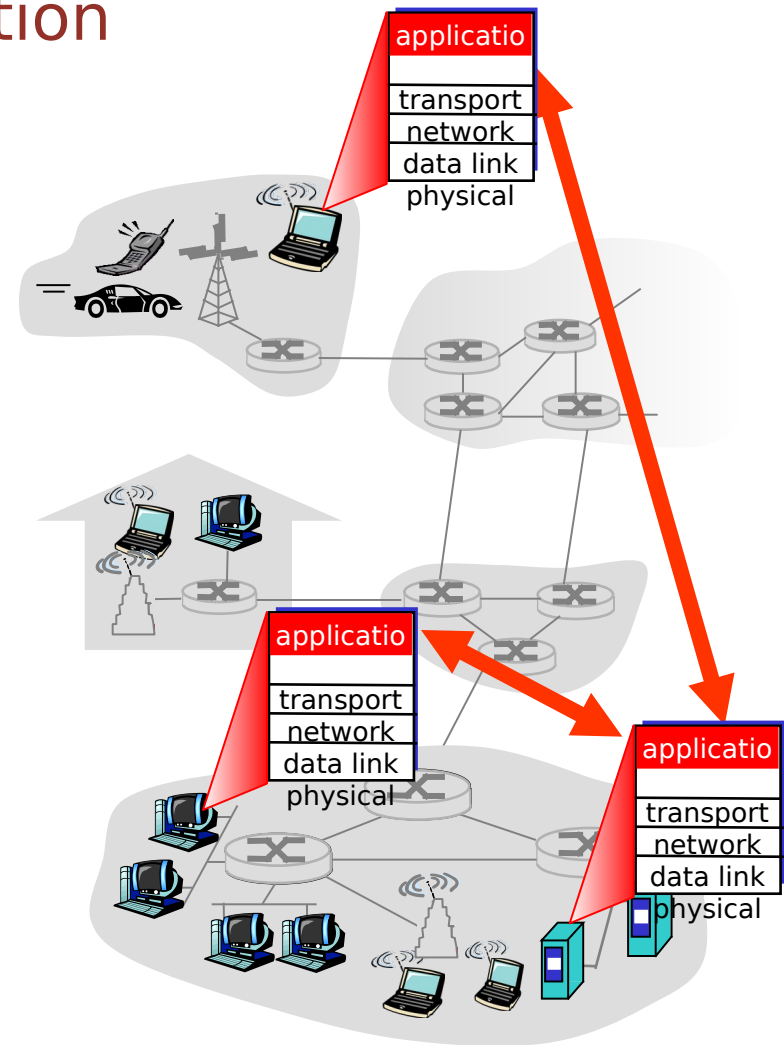
# Creating a network application

## write programs that

- run on (different) *end systems*
- communicate over network
- e.g., web server software communicates with browser software

## No need to write software for network-core devices

- network-core devices do not run user applications
- applications on end systems allows for rapid app development, propagation





## Some network apps

- e-mail
- remote login
- web
- instant messaging
- P2P file sharing
- multi-user network games
- streaming stored video (YouTube)
- voice over IP
- real-time video conferencing
- social networking
- cloud computing
- ...
- ...
- ...



# Network applications need streams of data

## Circuit switching —————> Packet switching



[http://www.tcpiguide.com/free/t\\_CircuitSwitchingandPacketSwitchingNetworks-2.htm](http://www.tcpiguide.com/free/t_CircuitSwitchingandPacketSwitchingNetworks-2.htm)

## Today's networks provide packet delivery, not streams!



Source: Freedman (partial)



## What if the Data Doesn't Fit?

```
GET /courses/archive/spr09/cos461/ HTTP/1.1  
Host: www.cs.princeton.edu  
User-Agent: Mozilla/4.03  
CRLF
```


Request



### Problem: Packet size

- Typical Web page is 10 kbytes
- On Ethernet, max IP packet is 1500 bytes

GET index.html



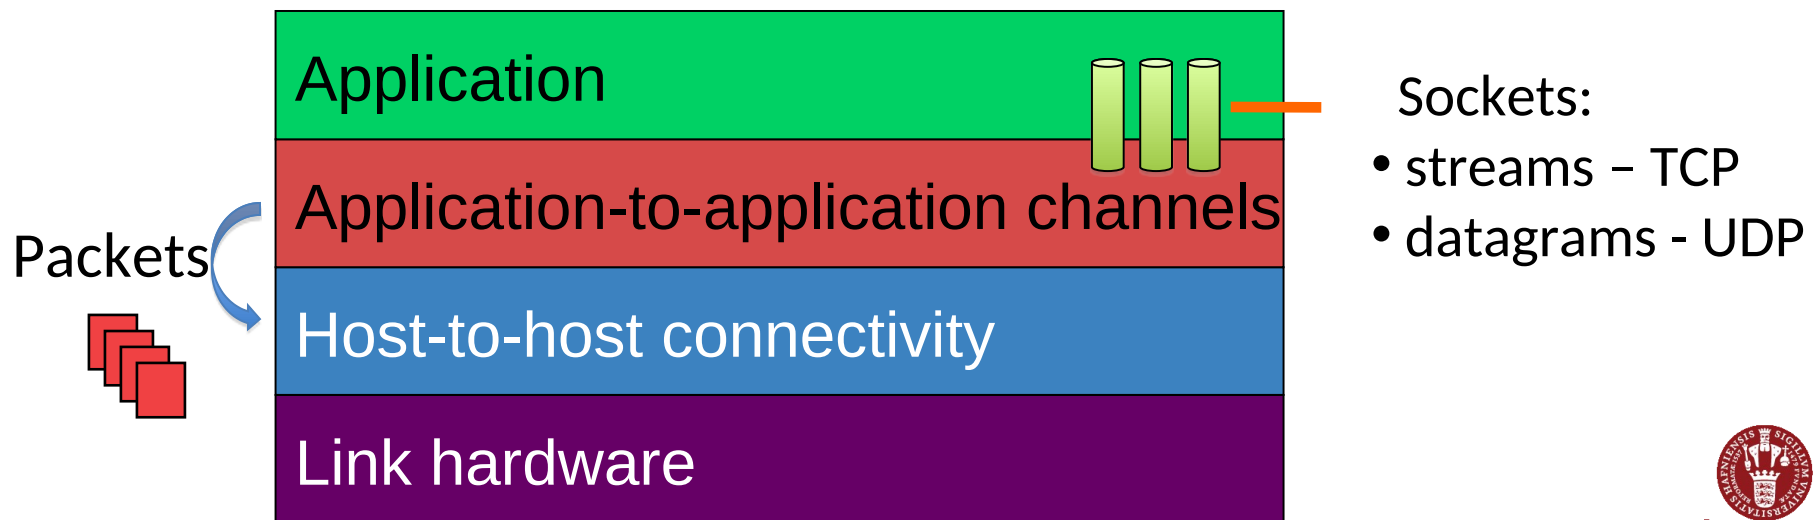
Solution: Split the data across multiple packets

Source: Freedman



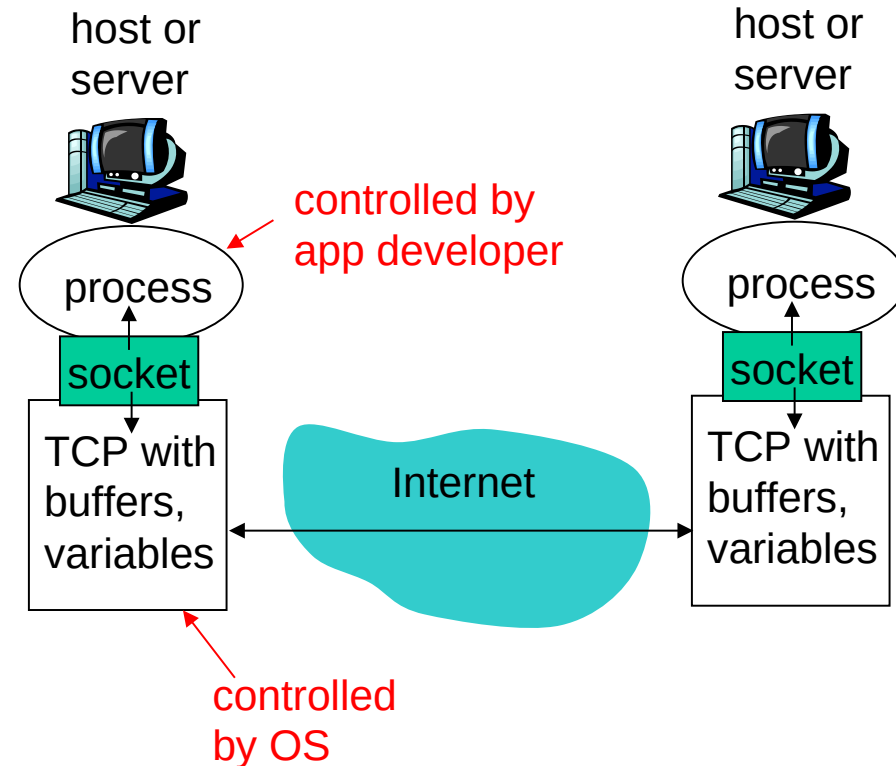
# Layering = Functional Abstraction

- Sub-divide the problem
  - Each layer relies on services from layer below
  - Each layer exports services to layer above
- Interface between layers defines interaction
  - Hides implementation details
  - Layers can change without disturbing other layers



# Sockets

- process sends/receives messages to/from its **socket**
- socket analogous to door
  - sending process shoves message out door
  - sending process relies on transport infrastructure on other side of door which brings message to socket at receiving process



API: (1) choice of transport protocol; (2) ability to fix a few parameters (more on this in next lecture!)

# Internet transport protocols services

## TCP service:

- *connection-oriented*: setup required between client and server processes
- *reliable transport*: between sending and receiving process
- *flow control*: sender won't overwhelm receiver
- *congestion control*: throttle sender when network overloaded
- *does not provide*: timing, minimum throughput guarantees, security

## UDP service:

- unreliable data transfer between sending and receiving process
- *does not provide*: connection setup, reliability, flow control, congestion control, timing, throughput guarantee, or security

Q: why bother? Why is there a UDP?



# UNIX Socket API

- Socket interface
  - Originally provided in Berkeley UNIX
  - Later adopted by all popular operating systems
  - Simplifies porting applications to different OSes
- In UNIX, everything is like a file
  - All input is like reading a file
  - All output is like writing a file
  - File is represented by an integer file descriptor
- API implemented as system calls
  - E.g., connect, read, write, close, ...

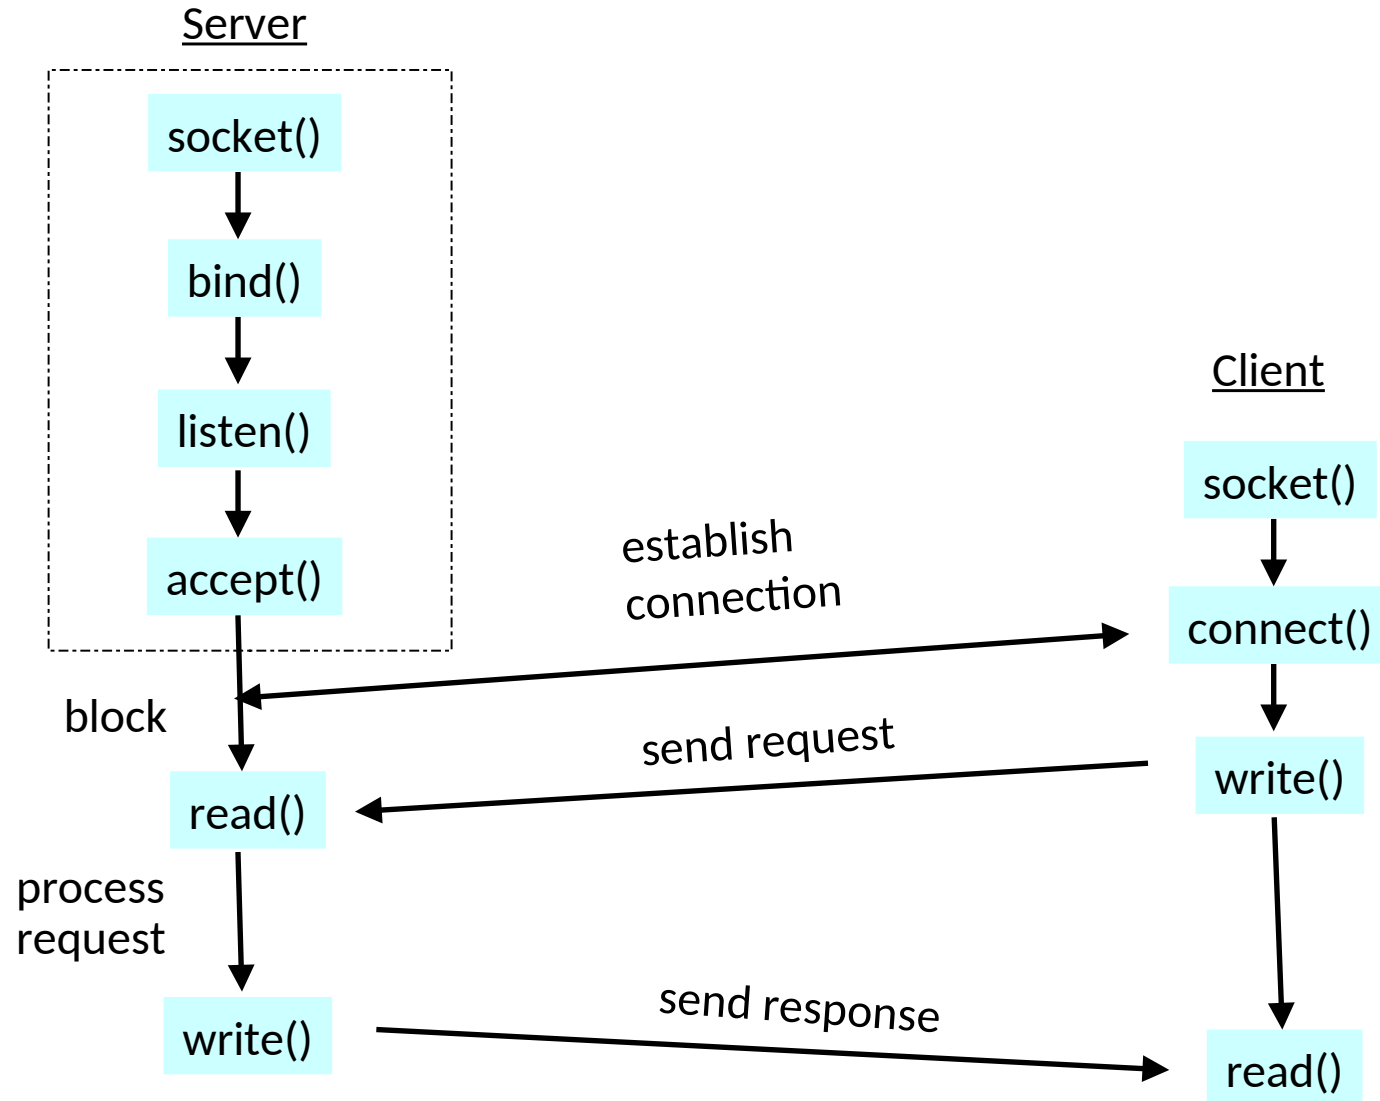


# Identifying the Receiving Process

- Sending process must identify the receiver
  - The receiving end host machine
  - The specific socket in a process on that machine
- Receiving host
  - Destination address that uniquely identifies the host
    - Typically, high-level name translated to IP address (DNS)
    - For example, [www.diku.dk](http://www.diku.dk) → **130.225.96.108**
  - An IP address is a 32-bit quantity
- Receiving socket
  - Host may be running many different processes
  - Destination port that uniquely identifies the socket
  - A port number is a 16-bit quantity




# Client-Server TCP Sockets



Source:  
Freedman  
(partial)




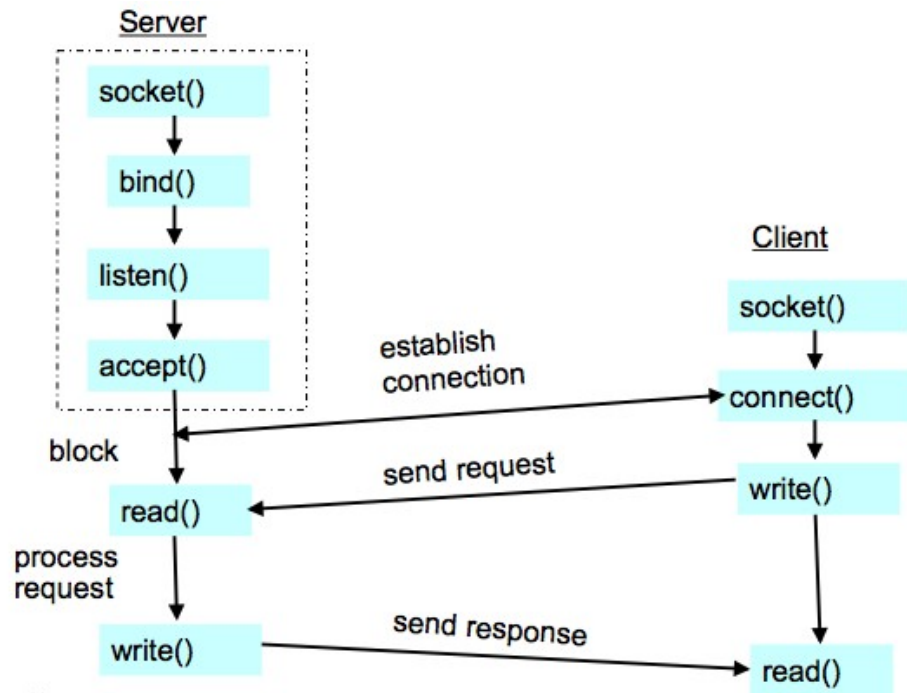
# Application Architectures



- P2P highly scalable, but difficult to manage
- Hybrids also possible, e.g., Skype



# Discussion: How do you set sockets up in a P2P application?



- How many sockets in each peer?
- What if many peers connect to one peer?

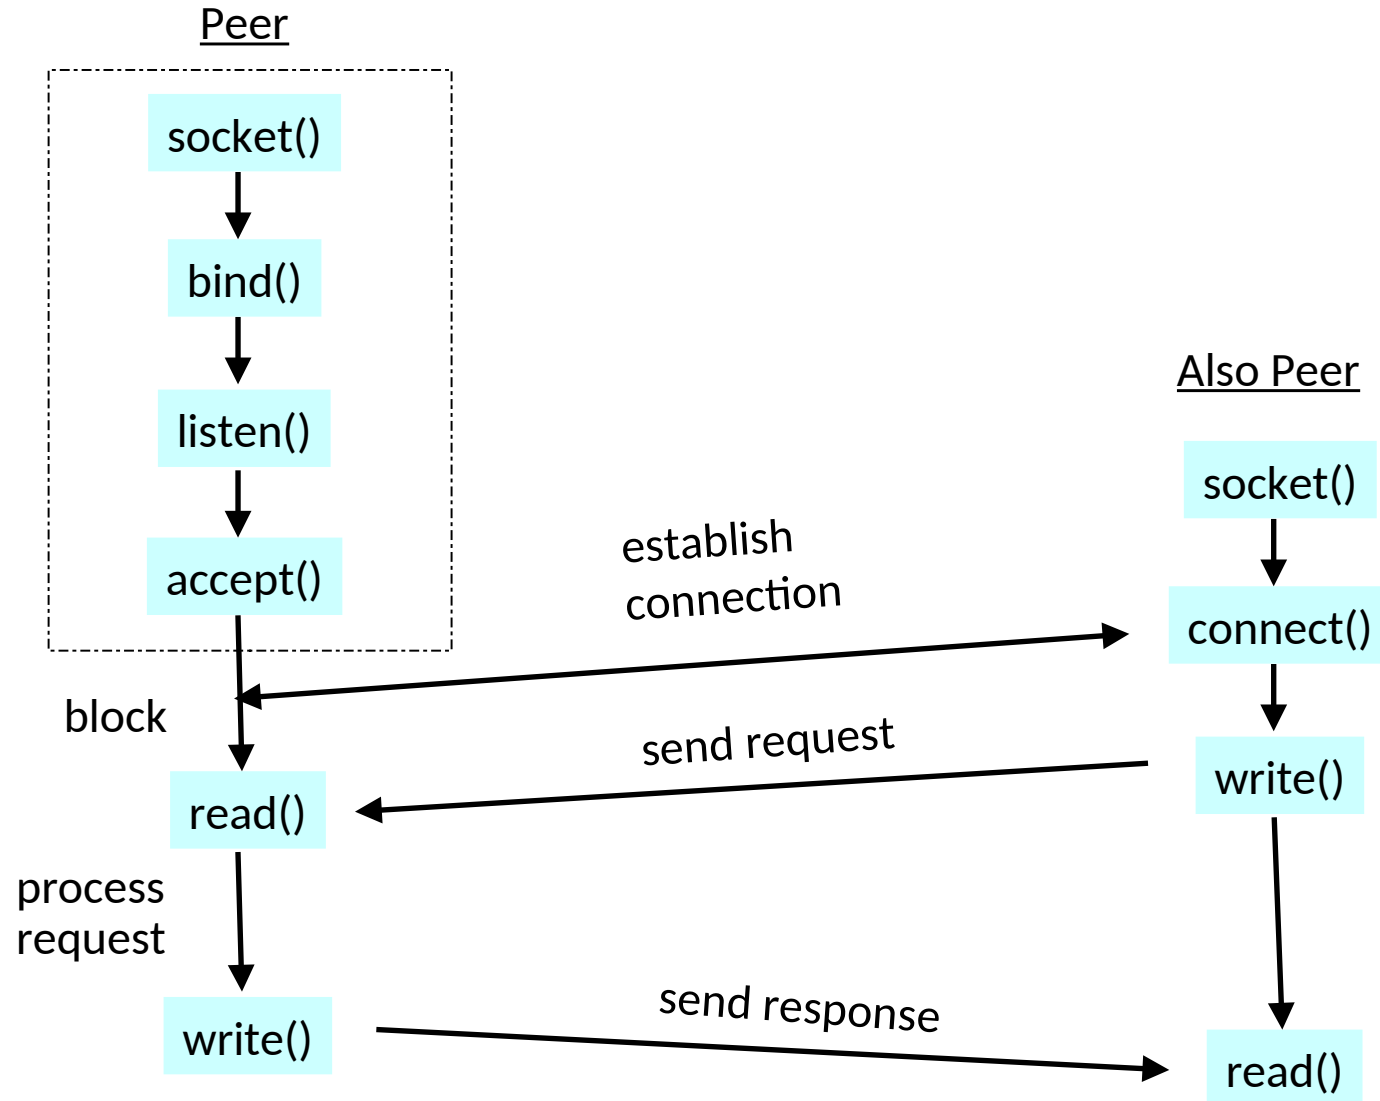
## P2P Sockets

- *The proof is trivial and left as an exercise to the reader.*

Carl-Johannes Johnsen



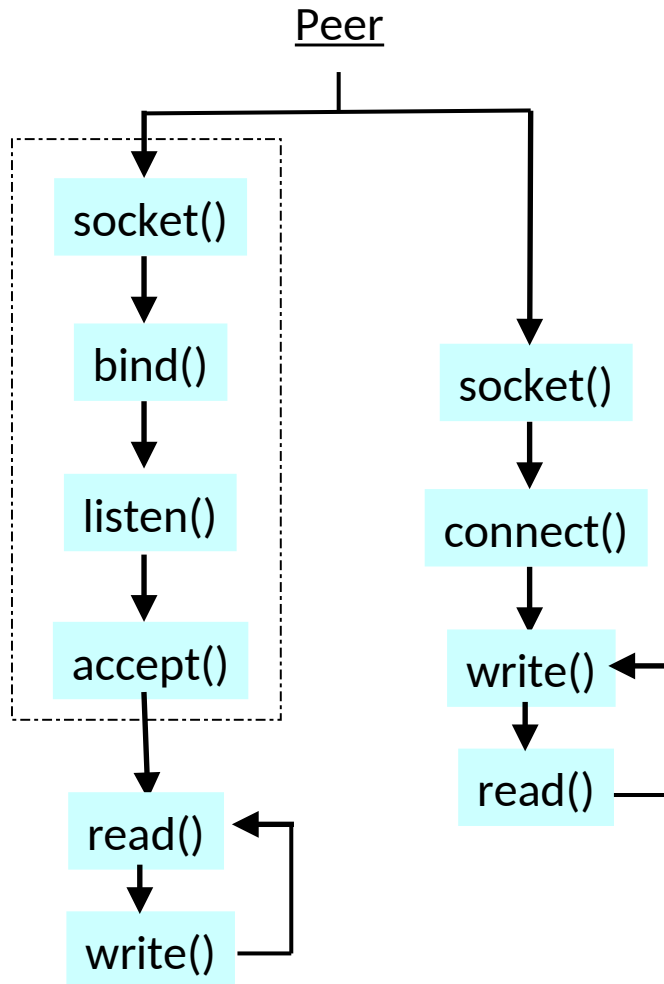
# P2P Sockets



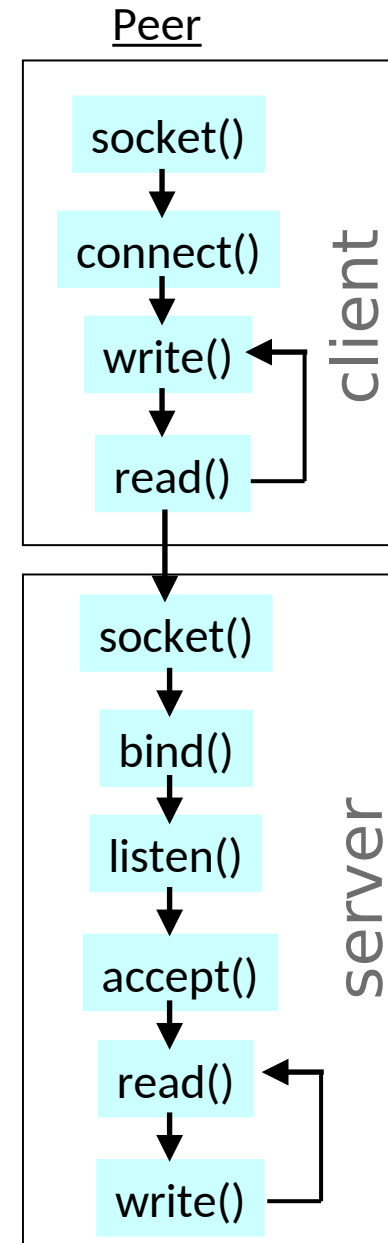
Source:  
Freedman  
(partial)



# P2P Sockets



OR




Source:  
Freedman  
(partial)



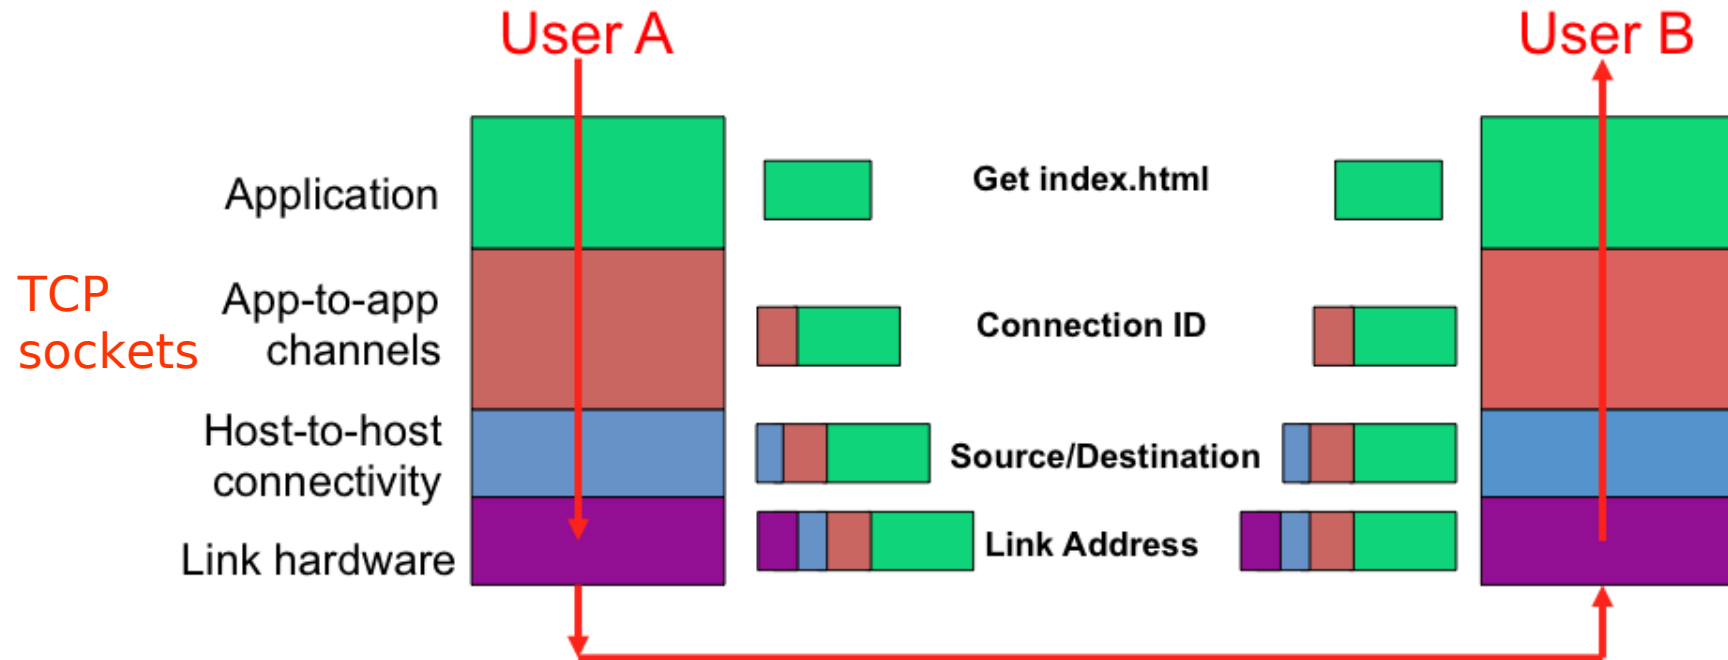
# HTTP Basics

- HTTP layered over bidirectional byte stream
- Interaction
  - Client sends request to server, followed by response from server to client
  - Requests/responses are encoded in text
- Targets access to web objects
  - GET, POST, HEAD → HTTP/1.0
  - GET, POST, HEAD, PUT, DELETE → HTTP/1.1
- Stateless
  - Server maintains no info about past client requests
    - What about personalization? Data stored in back-end database; client sends “ web cookie” used to lookup data





## Layer Encapsulation in HTTP



# HTTP Request Example

GET / HTTP/1.1

Host: sns.cs.princeton.edu

Accept: \*/\*

Accept-Language: en-us

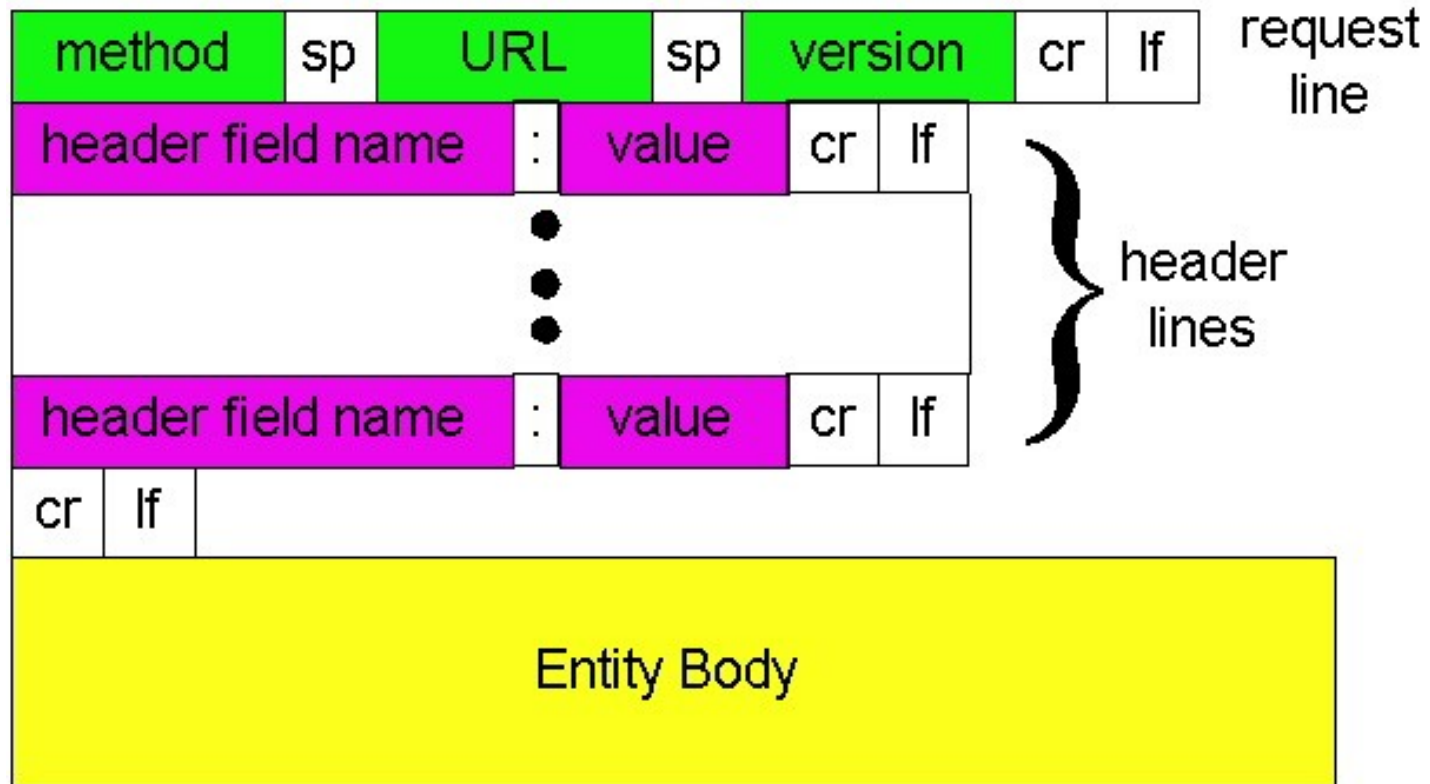
Accept-Encoding: gzip, deflate

User-Agent: Mozilla/5.0 (Macintosh; U; Intel Mac OS X 10.5; en-US; rv:1.9.2.13) Gecko/20101203 Firefox/3.6.13

Connection: Keep-Alive



# HTTP Request





# HTTP Request

GET / HTTP/1.1

Host: sns.cs.princeton.edu

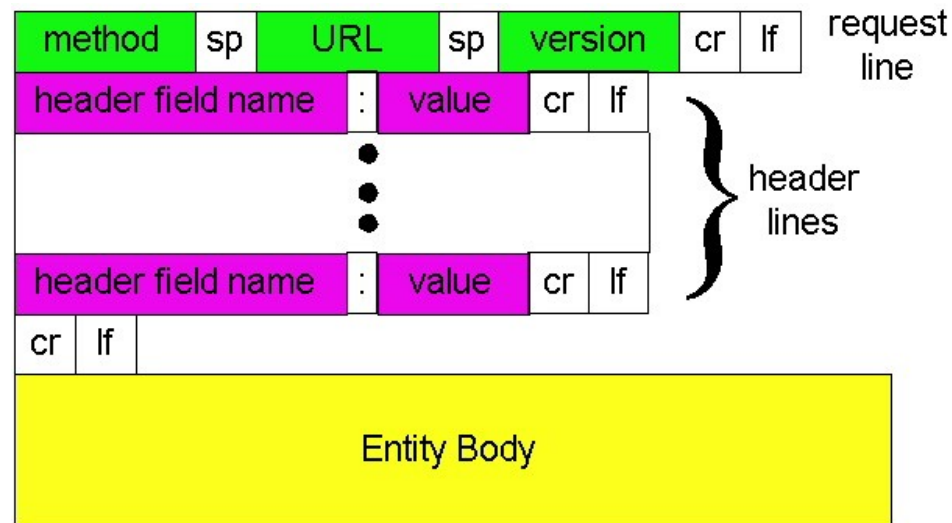
Accept: \*/\*

Accept-Language: en-us

Accept-Encoding: gzip, deflate

User-Agent: Mozilla/5.0 (Macintosh; U; Intel Mac OS X 10.5; en-US; rv:1.9.2.13)  
Gecko/20101203 Firefox/3.6.13

Connection: Keep-Alive



# HTTP Response Example

HTTP/1.1 200 OK

Date: Wed, 02 Feb 2011 04:01:21 GMT

Server: Apache/2.2.3 (CentOS)

X-Pingback: <http://sns.cs.princeton.edu/xmlrpc.php>

Last-Modified: Wed, 01 Feb 2011 12:41:51 GMT

ETag: "7a11f-10ed-3a75ae4a"

Accept-Ranges: bytes

Content-Length: 4333


Keep-Alive: timeout=15, max=100

Connection: Keep-Alive

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Transitional//EN"
    "http://www.w3.org/TR/xhtml1/DTD/xhtml1-transitional.dtd">
<html xmlns="http://www.w3.org/1999/xhtml" dir="ltr" lang="en-US">
```



# HTTP Response




## Some HTTP Response codes

200	OK
204	No Content
301	Moved Permanently
400	Bad Request
403	Forbidden
404	Not Found
418	I'm a teapot
500	Internal Server Error



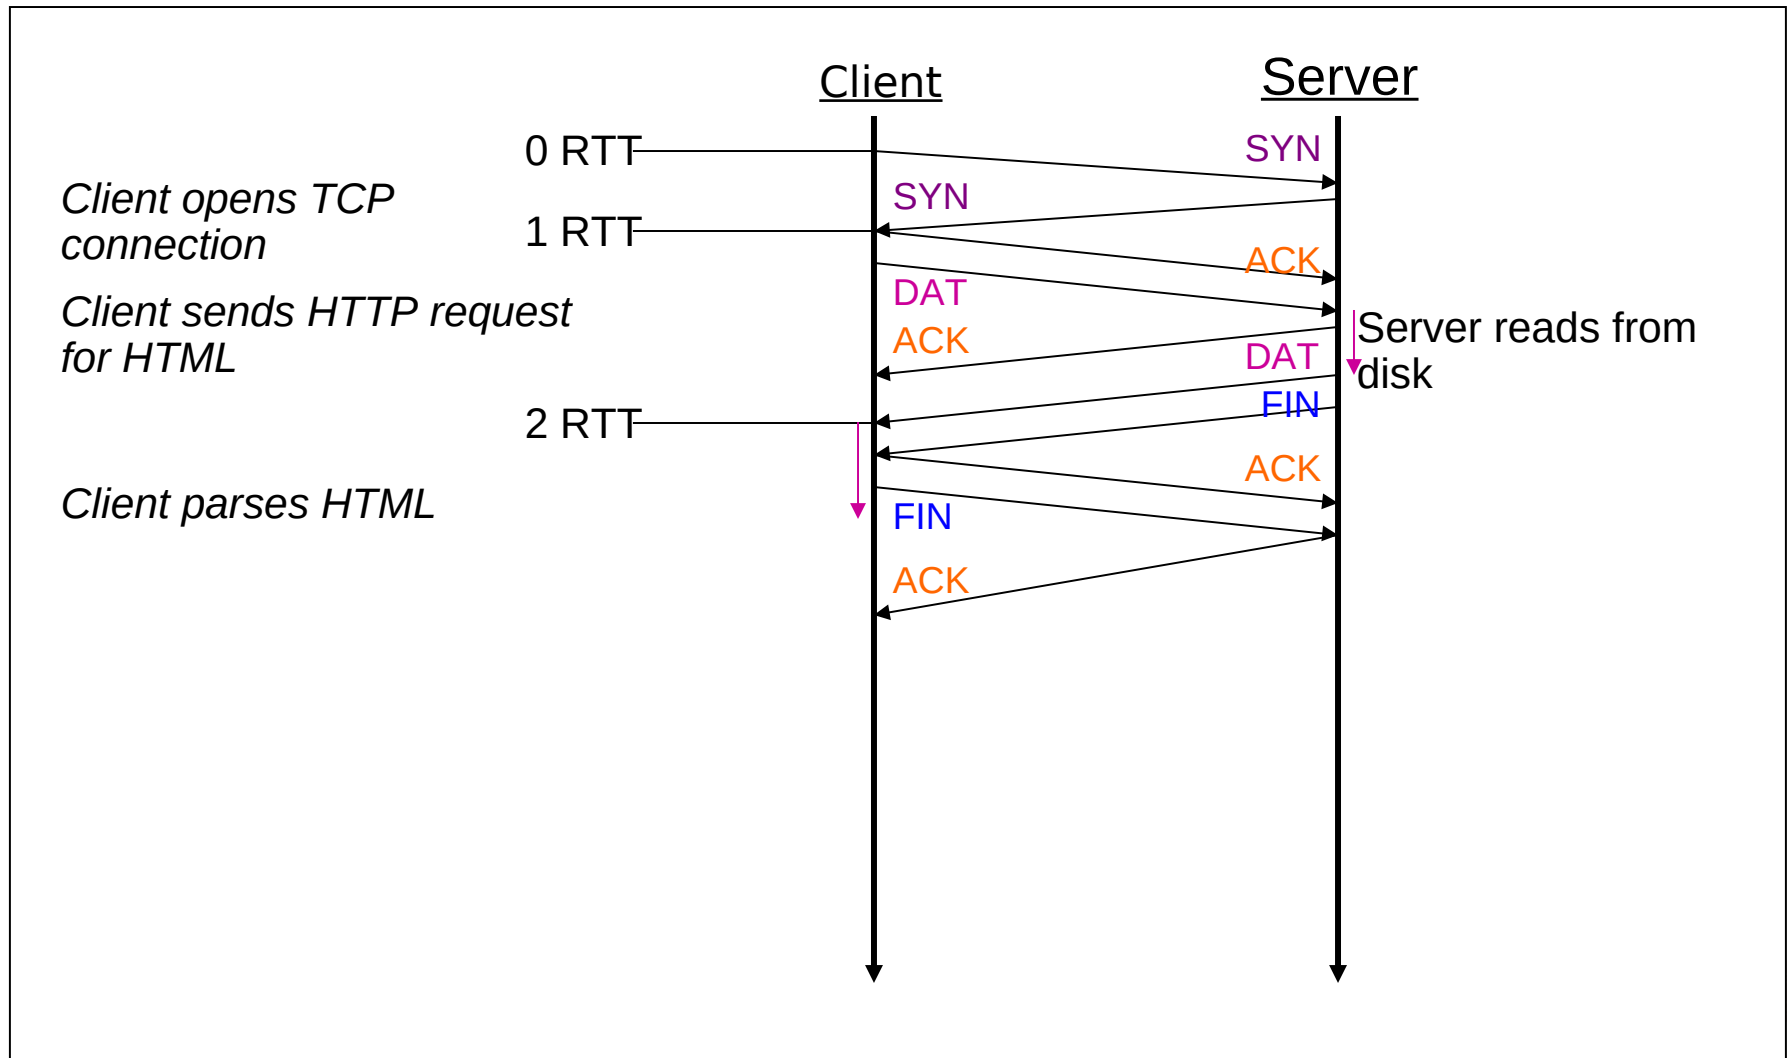
# Single Transfer Example

Source: Freedman



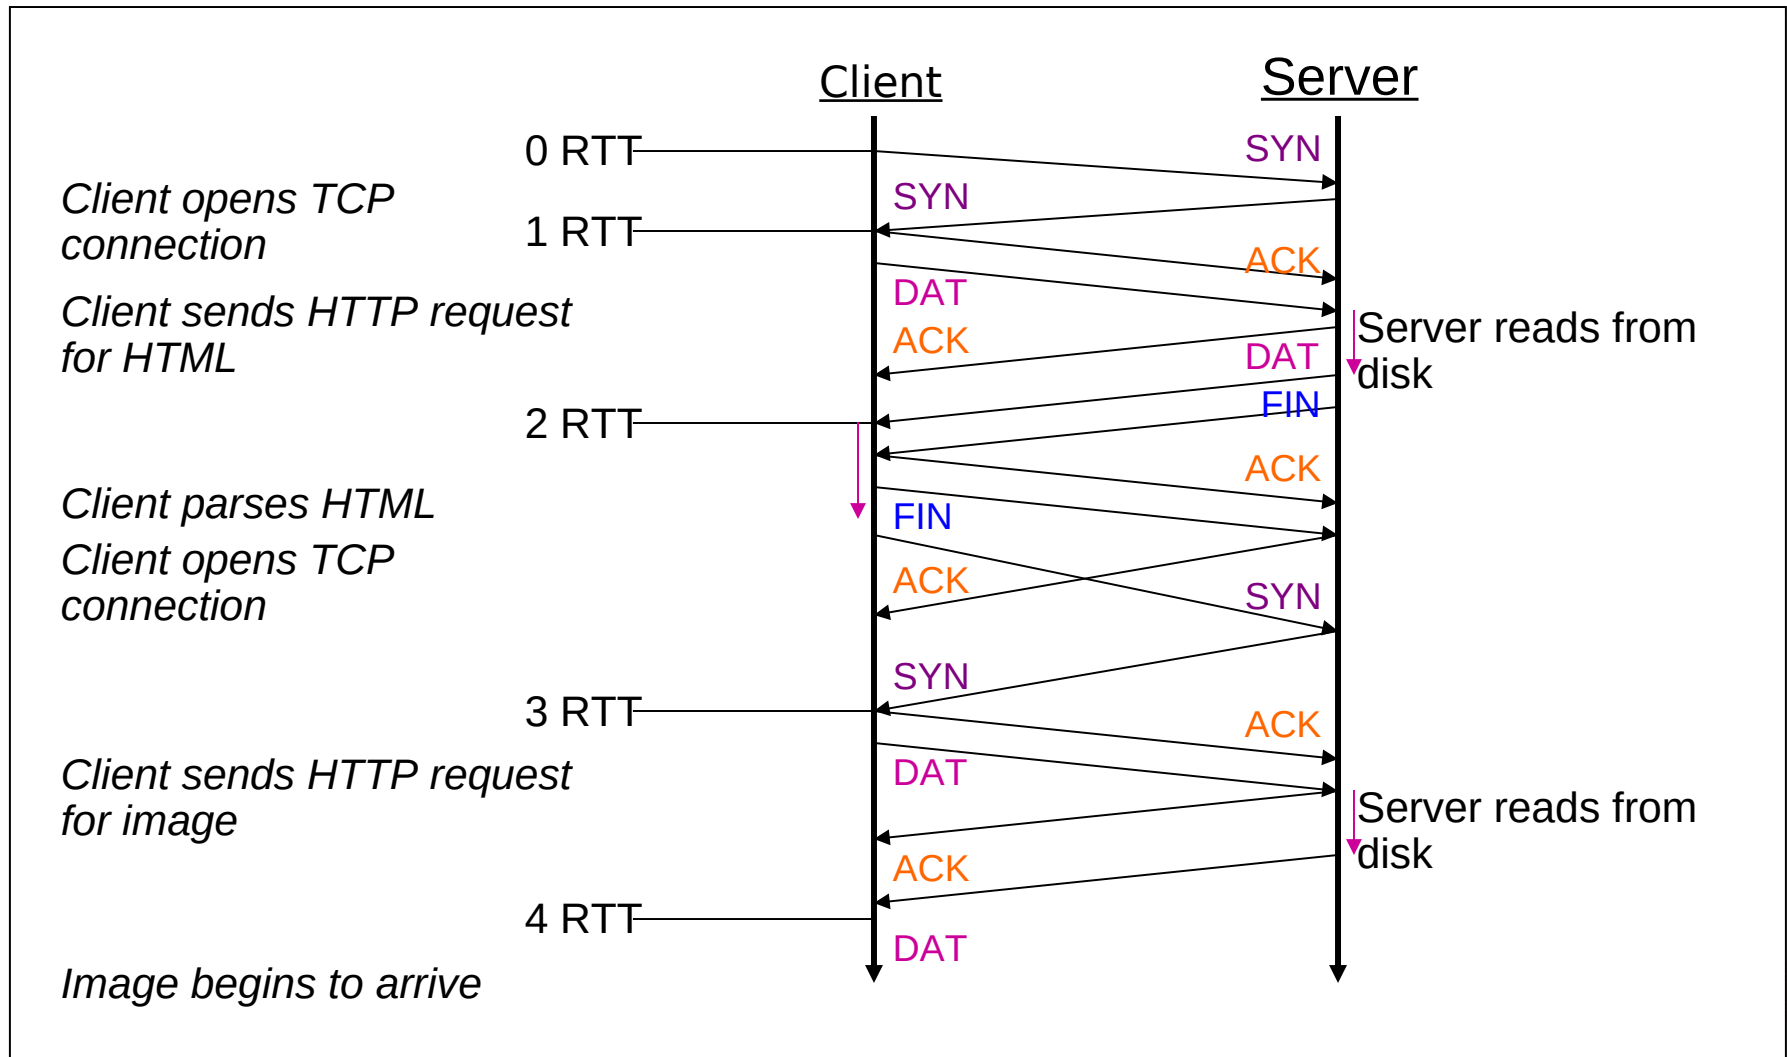
# Single Transfer Example

Source: Freedman



# Single Transfer Example

Source: Freedman



## Problems with simple model

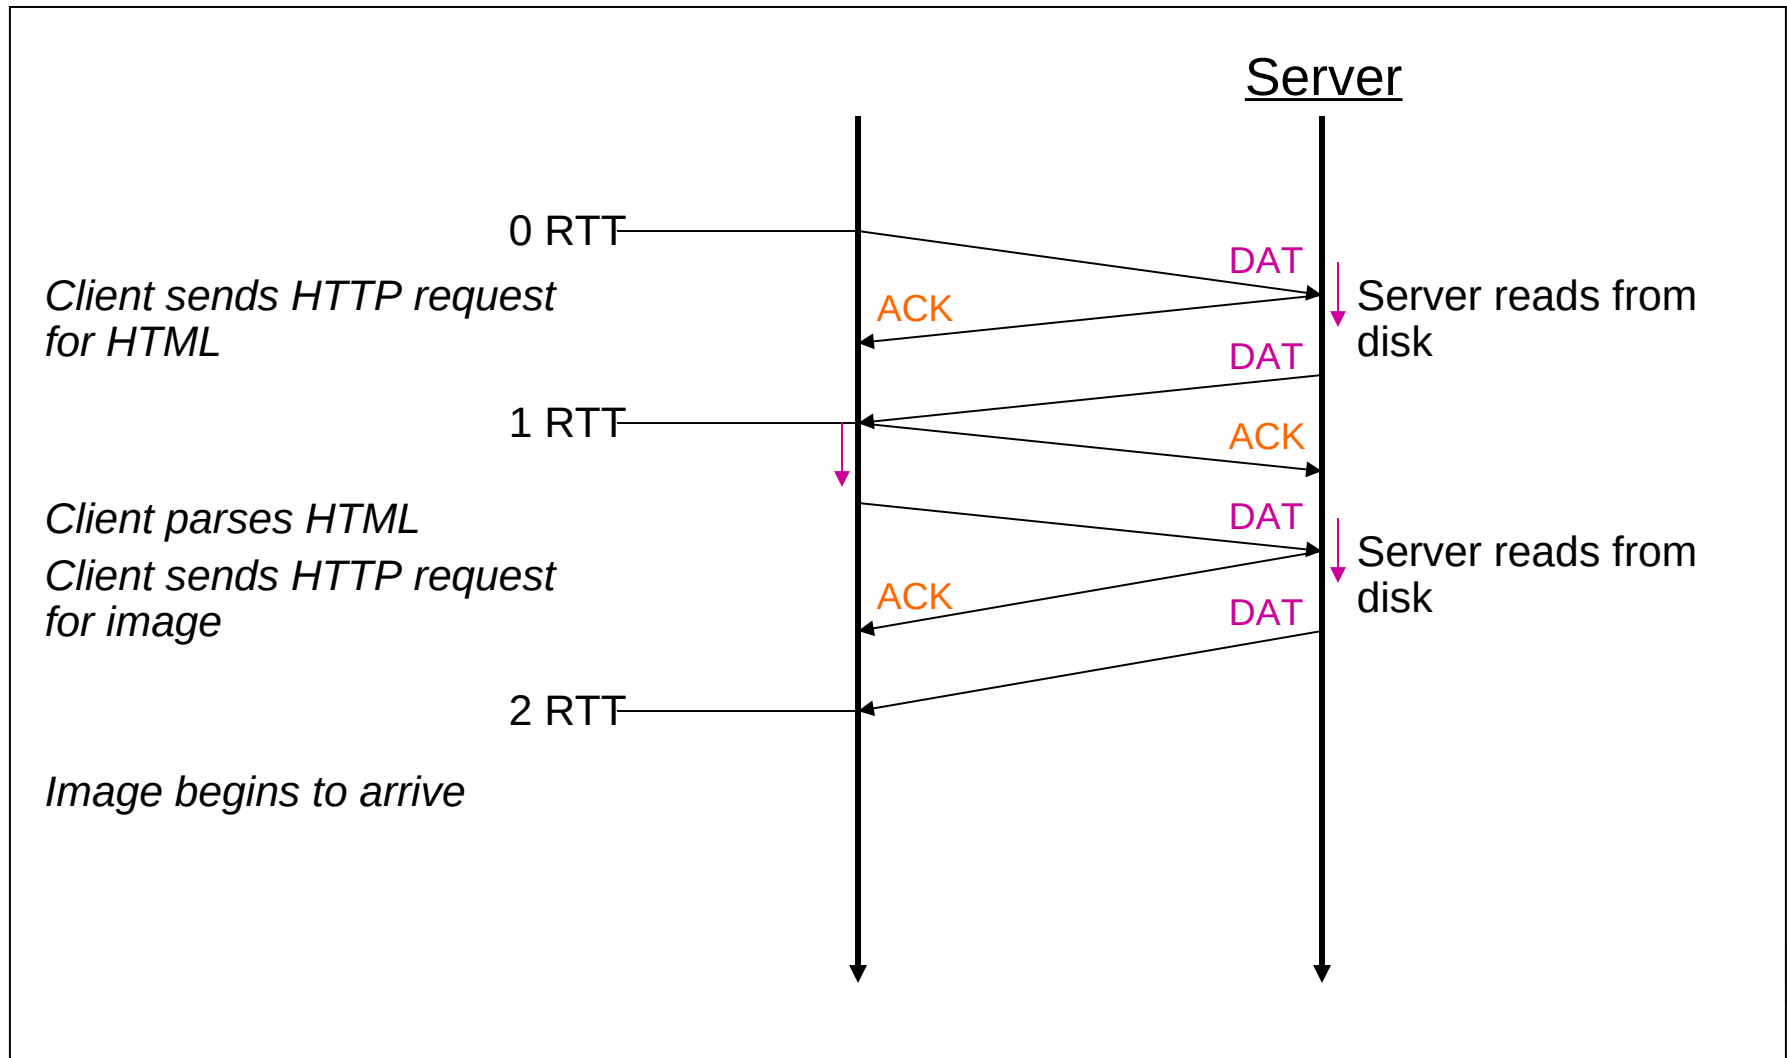
- Multiple connection setups
  - Three-way handshake each time (TCP “synchronizing” stream)
- Lots of extra connections
  - Increases server state/processing
  - Server forced to keep connection state
- Later we will see also that
  - Short transfers are hard on stream protocol (TCP)
    - How much data should it send at once?
    - Congestion avoidance: Takes a while to “ramp up” to high sending rate (TCP “slow start”)
    - Loss recovery is poor when not “ramped up”





# Persistent Connection Example

Source: Freedman



# Persistent HTTP

## Non-persistent HTTP

### issues:

- Requires 2 RTTs per object
- OS must allocate resources for each TCP connection
- But browsers often open parallel TCP connections to fetch referenced objects

## Persistent HTTP:

- Server leaves connection open after sending response
- Subsequent HTTP messages between same client/server are sent over connection



# Persistent HTTP

## Persistent without pipelining:

- Client issues new request only when previous response has been received
- One RTT for each object

## Persistent with pipelining:

- Default in HTTP/1.1 spec
- Client sends requests as soon as it encounters referenced object
- As little as one RTT for all the referenced objects
- Server must handle responses in same order as requests

• Persistent without pipelining most common:  
When does pipelining work best?

• Multiple parallel requests or pipelined requests?

Source: Freedman (partial)



# HTTP Caching

- Clients often cache documents
  - When should origin be checked for changes?
  - Every time? Every session? Date?
- HTTP includes caching information in headers
  - HTTP 0.9/1.0 used: “ Expires: <date>” ; “ Pragma: no-cache”
  - HTTP/1.1 has “ Cache-Control”
    - “ No-Cache” , “ Private” , “ Max-age: <seconds>”
    - “ E-tag: <opaque value>”
- If not expired, use cached copy
- If expired, use condition GET request to origin
  - “ If-Modified-Since: <date>” , “ If-None-Match: <etag>”
  - 304 ( “ Not Modified” ) or 200 ( “ OK” ) response



# HTTP Conditional Request

GET / HTTP/1.1

Host: sns.cs.princeton.edu

User-Agent: Mozilla/5.0 (Macintosh; U; Intel Mac OS X  
10.5; en-US; rv:1.9.2.13)

Connection: Keep-Alive

If-Modified-Since: Tue, 1 Feb 2011 17:54:18 GMT

If-None-Match: "7a11f-10ed-3a75ae4a"

HTTP/1.1 304 Not Modified

Date: Wed, 02 Feb 2011 04:01:21  
GMT

Server: Apache/2.2.3 (CentOS)

ETag: "7a11f-10ed-3a75ae4a"

Accept-Ranges: bytes


Keep-Alive: timeout=15, max=100

Connection: Keep-Alive




# Web Proxy Caches

- User configures browser: Web accesses via cache
- Browser sends all HTTP requests to cache
  - Object in cache: cache returns object
  - Else: cache requests object from origin, then returns to client



## When a single cache isn't enough

- What if the working set is  $>$  proxy disk?
  - Cooperation!
- A static hierarchy
  - Check local
  - If miss, check siblings
  - If miss, fetch through parent
- Internet Cache Protocol (ICP)
  - ICPv2 in RFC 2186 (& 2187)
  - UDP-based, short timeout



# Streaming multimedia: DASH

- *DASH: D*ynamic, *A*daptive *S*treaming over *H*TTP
- *server:*
  - divides video file into multiple chunks
  - each chunk stored, encoded at different rates
  - *manifest file:* provides URLs for different chunks
- *client:*
  - periodically measures server-to-client throughput
  - consulting manifest, requests one chunk at a time
    - chooses maximum coding rate sustainable given current throughput
    - can choose different coding rates at different points in time (depending on available bandwidth at time)





# Streaming multimedia: DASH

- *DASH: Dynamic, Adaptive Streaming over HTTP*
- “*intelligence*” at client: client determines
  - *when* to request chunk (so that buffer starvation, or overflow does not occur)
  - *what encoding rate* to request (higher quality when more bandwidth available)
  - *where* to request chunk (can request from URL server that is “close” to client or has high available bandwidth)

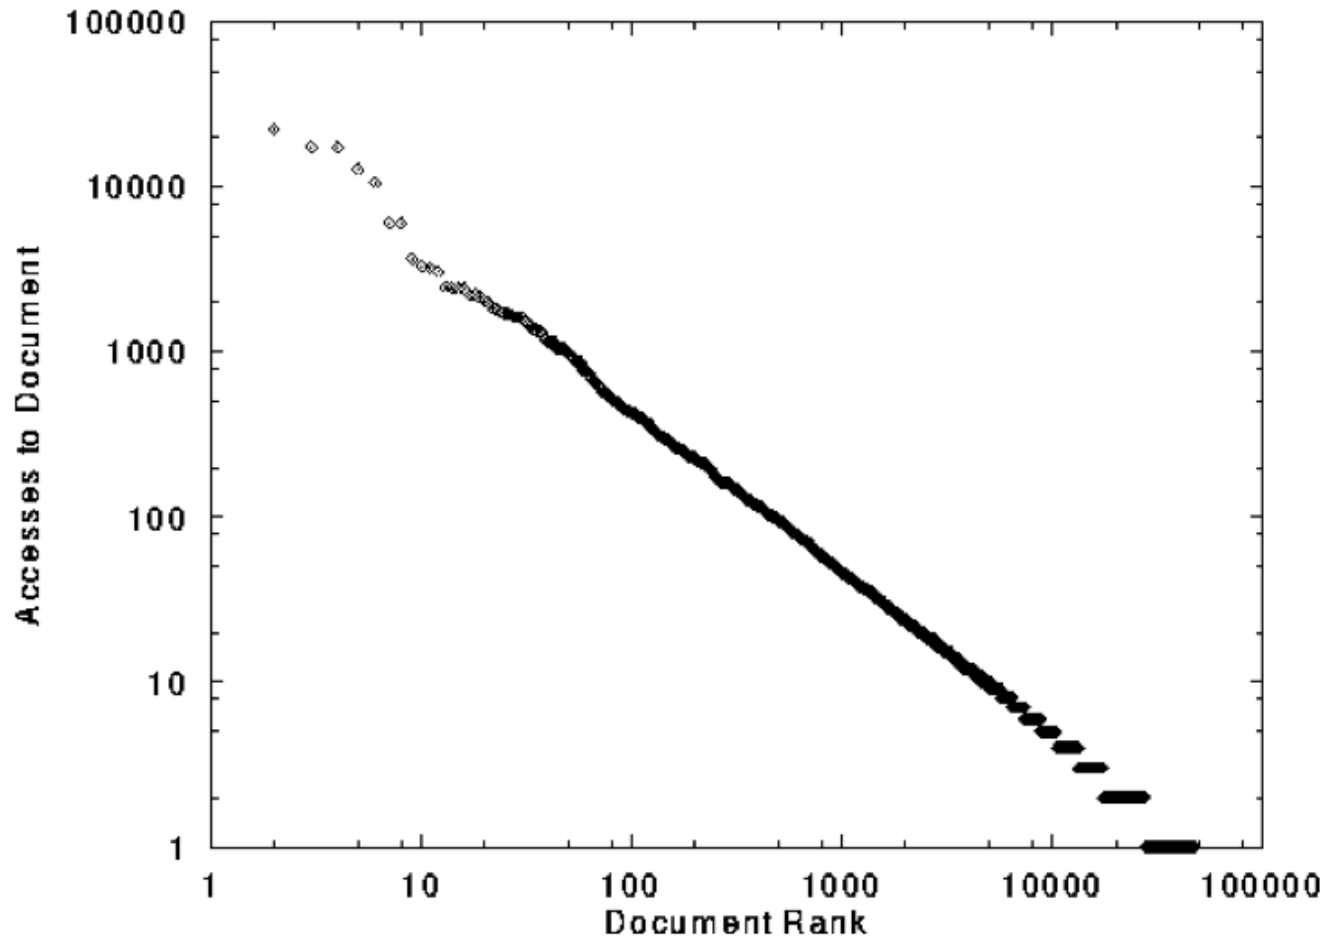


## Distributing Data to Clients

- Data needs to be close to clients
- Sending repeated data needs to be minimised
- Points of failure need to be minimised



## Web traffic has cacheable workload



“ Zipf” or “ power-law” distribution

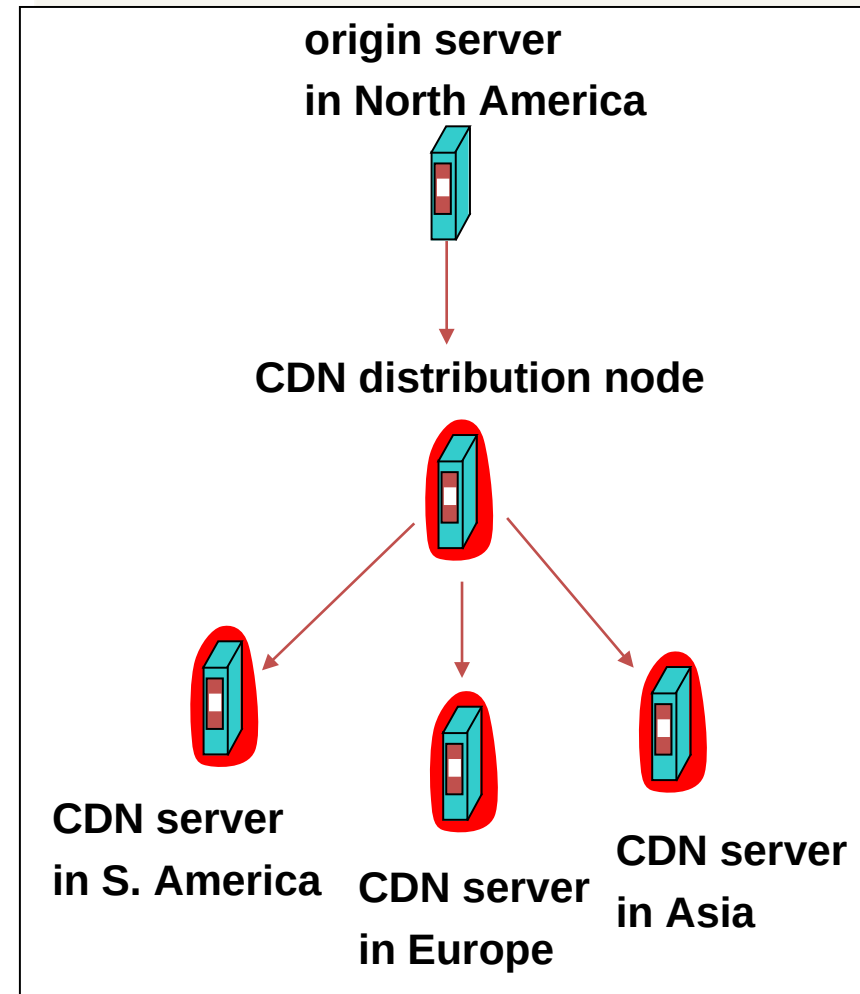
Source: Freedman

# Content Distribution Networks (CDNs)

- Content providers are CDN customers

## Content replication

- CDN company installs thousands of servers throughout Internet
  - In large datacenters (Bring Home)
  - Or, close to users (Enter Deep)
- CDN replicates customers' content
- When provider updates content, CDN updates servers



# Content Distribution Networks & Server Selection

- Replicate content on many servers
- Challenges
  - How to replicate content
  - Where to replicate content
  - How to find replicated content
  - How to choose among known replicas
  - How to direct clients towards replica



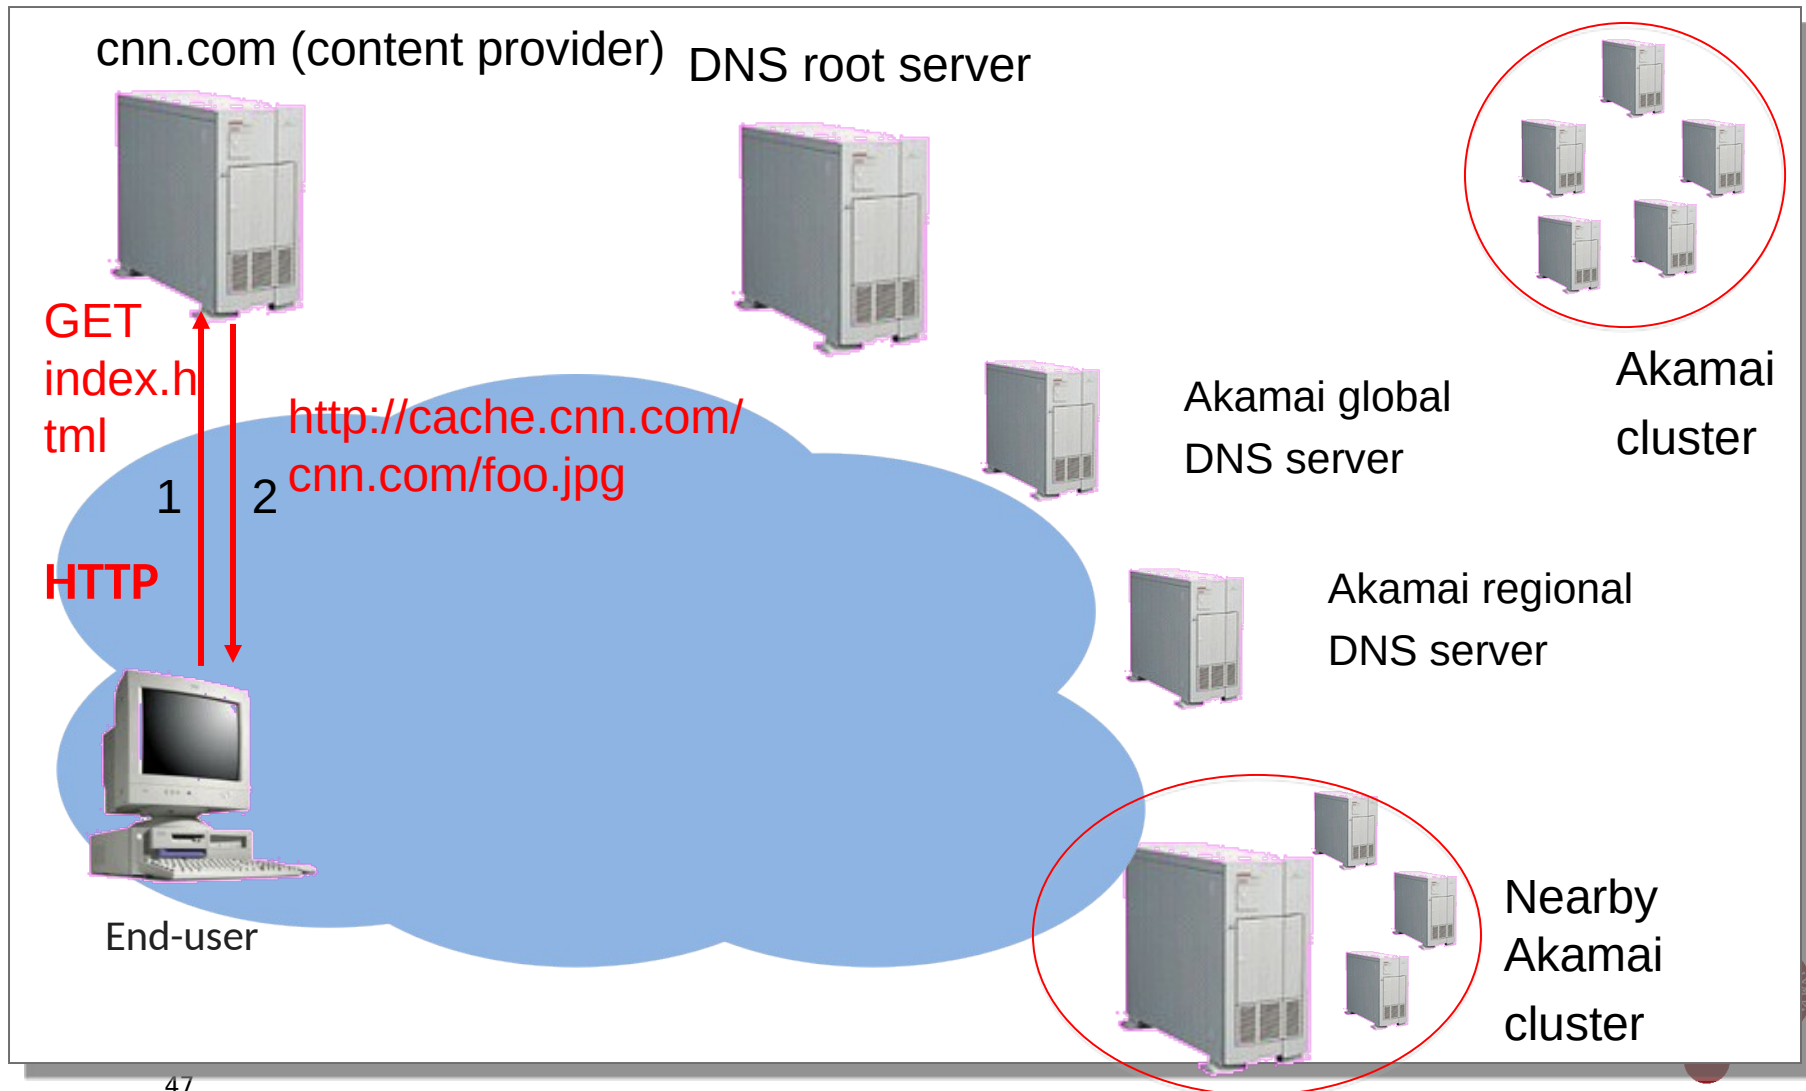
# Server Selection

- Which server?
  - Lowest load: to balance load on servers
  - Best performance: to improve client performance
    - Based on Geography? RTT? Throughput? Load?
  - Any alive node: to provide fault tolerance
- How to direct clients to a particular server?
  - As part of routing: anycast, cluster load balancing
  - As part of application: HTTP redirect
  - As part of naming: DNS
- We will explain some of these techniques better later in the course!



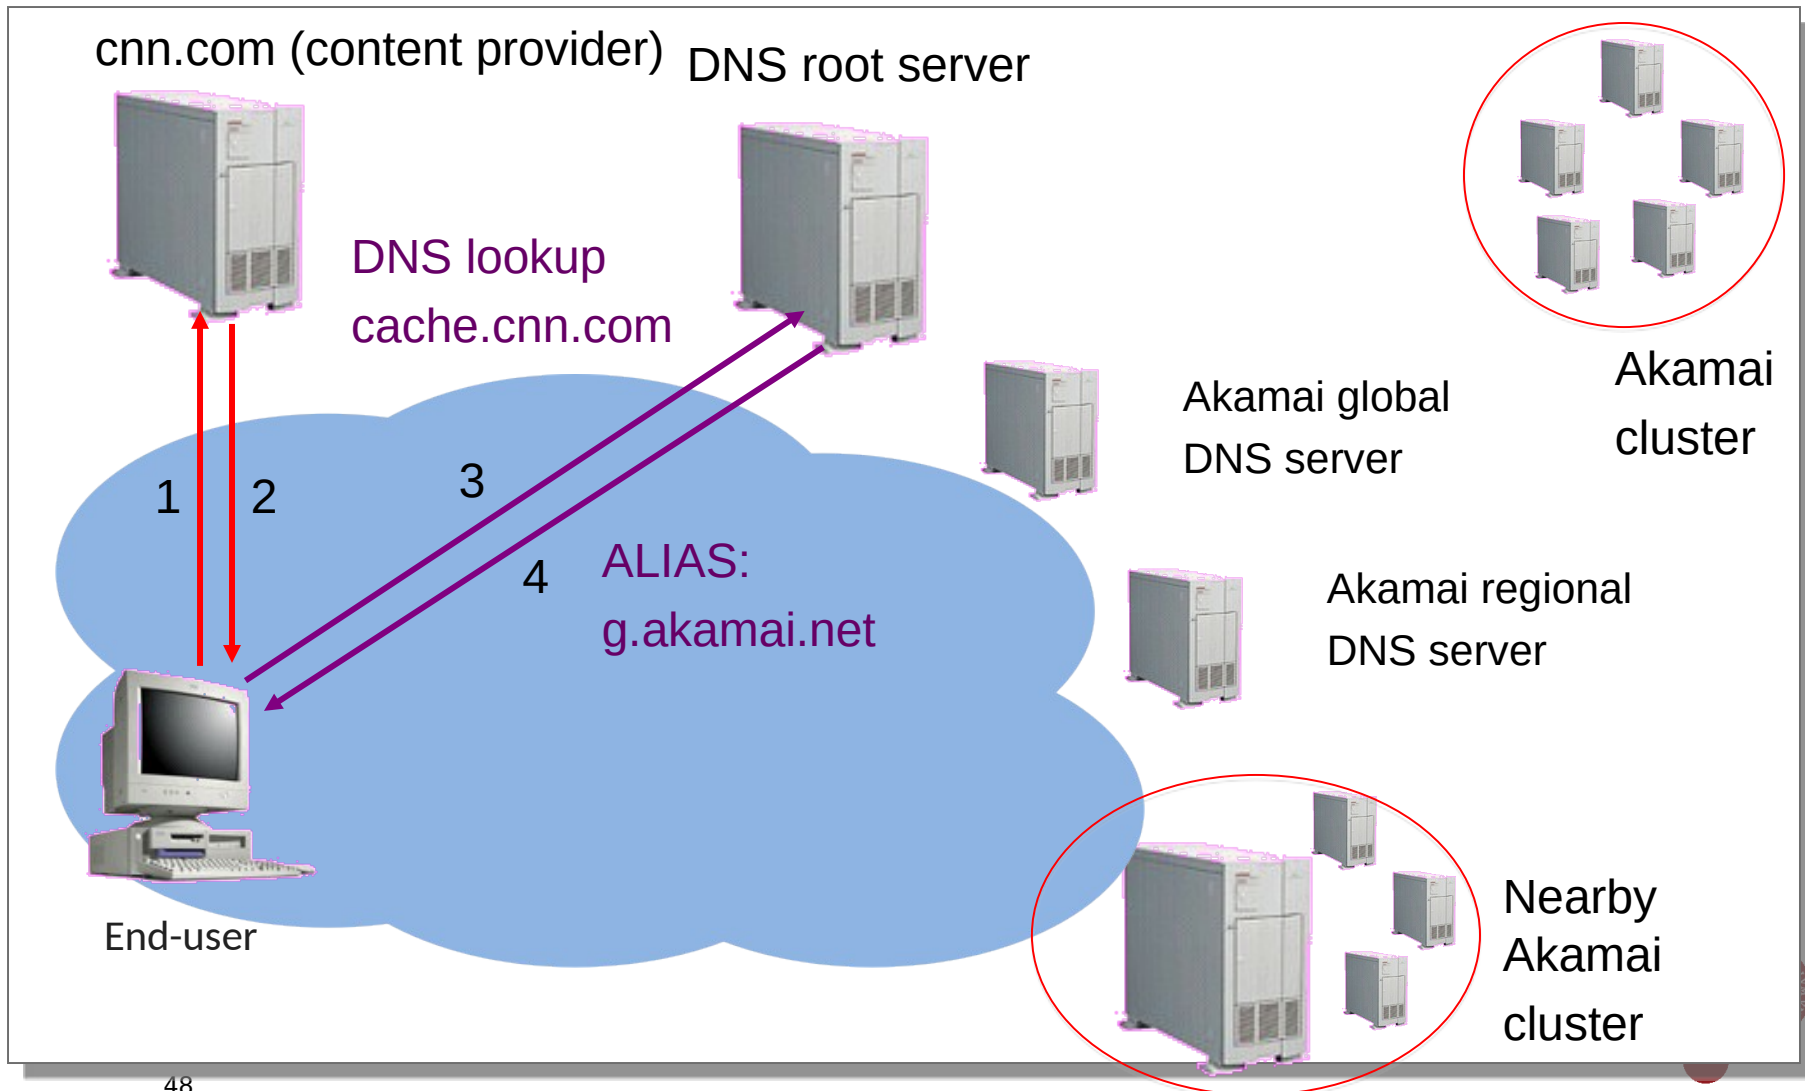
# How Akamai Works

Source: Freedman



# How Akamai Works

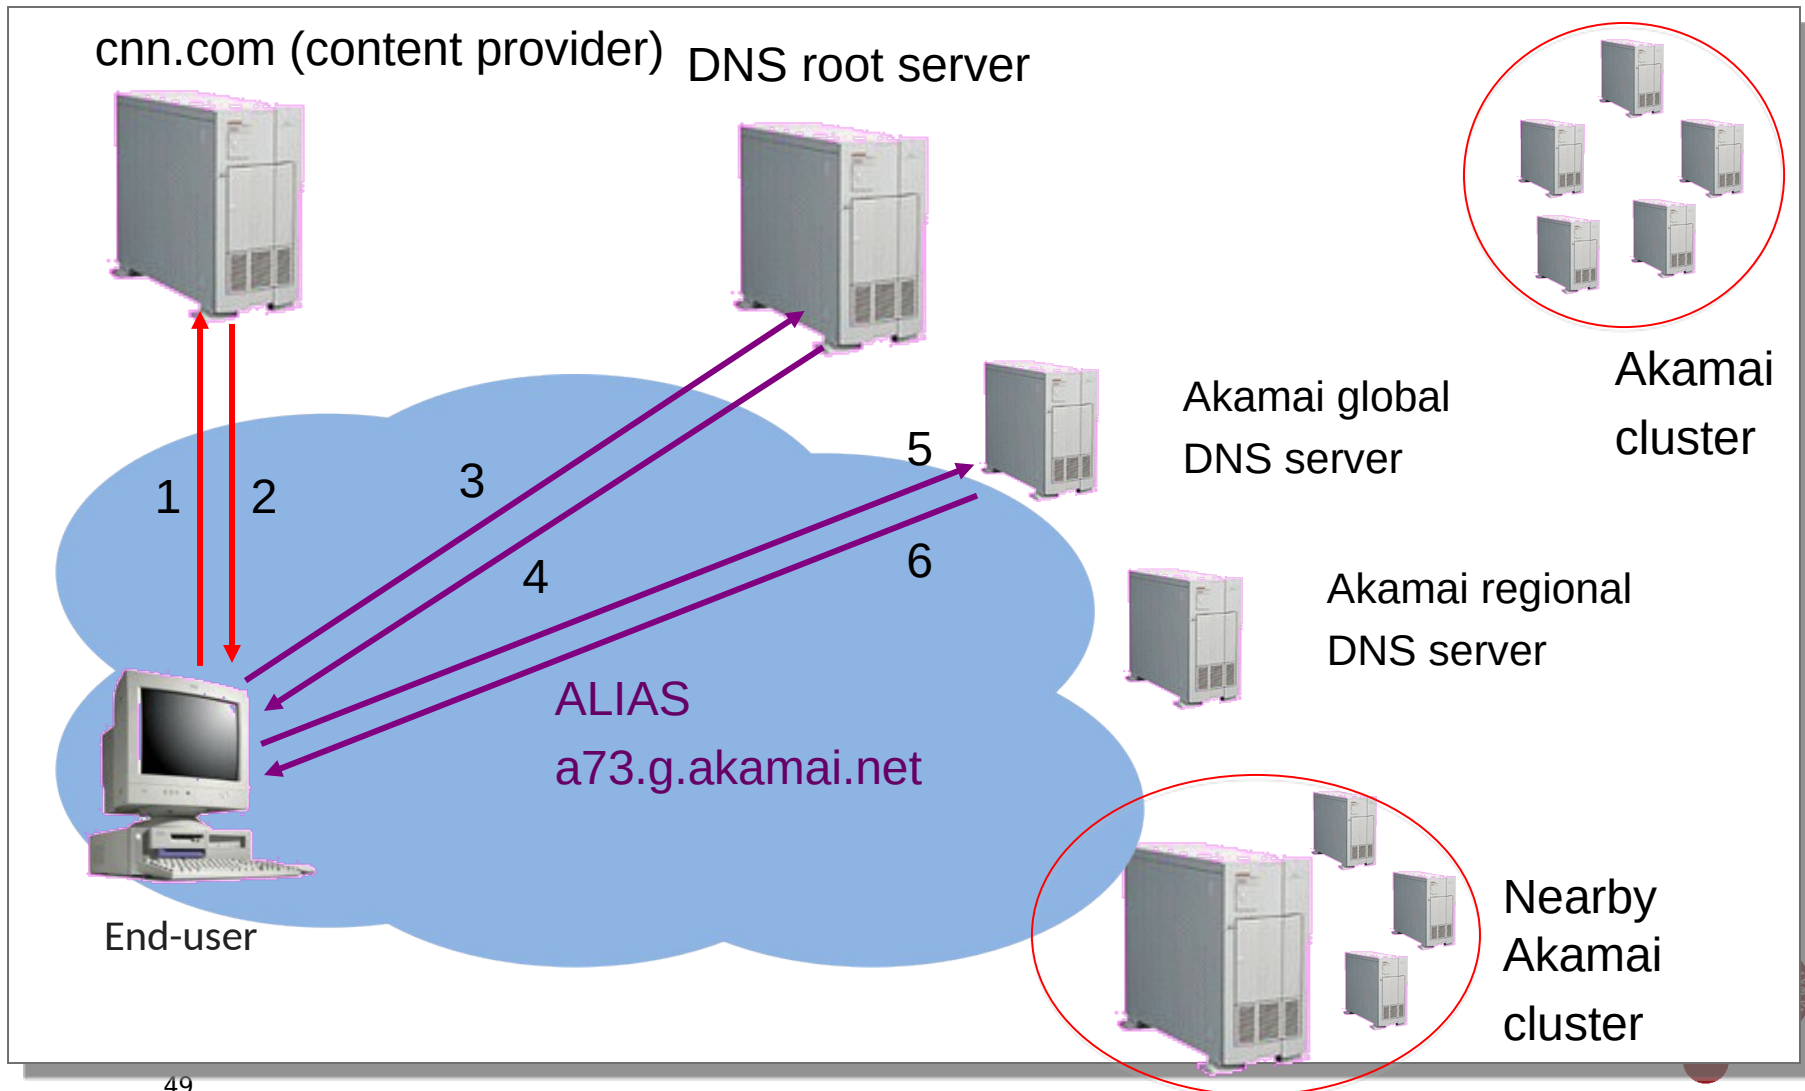
Source: Freedman





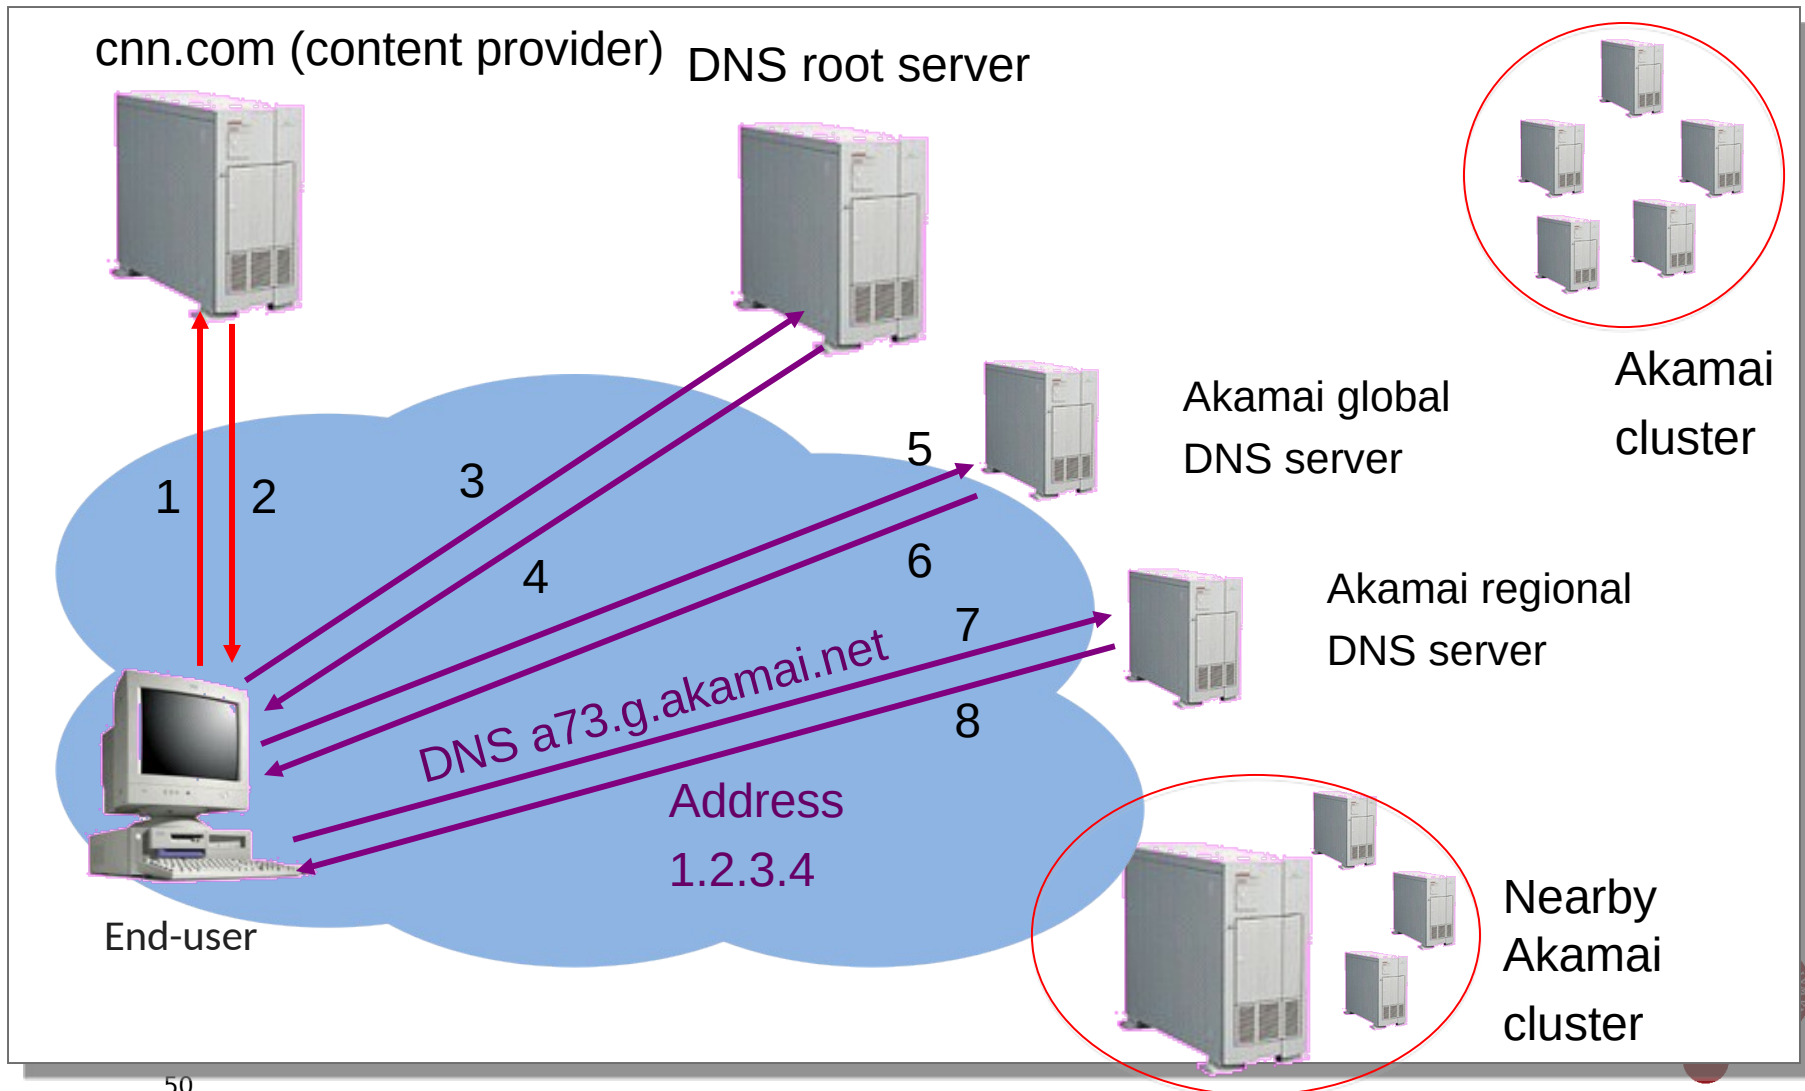
# How Akamai Works

Source: Freedman



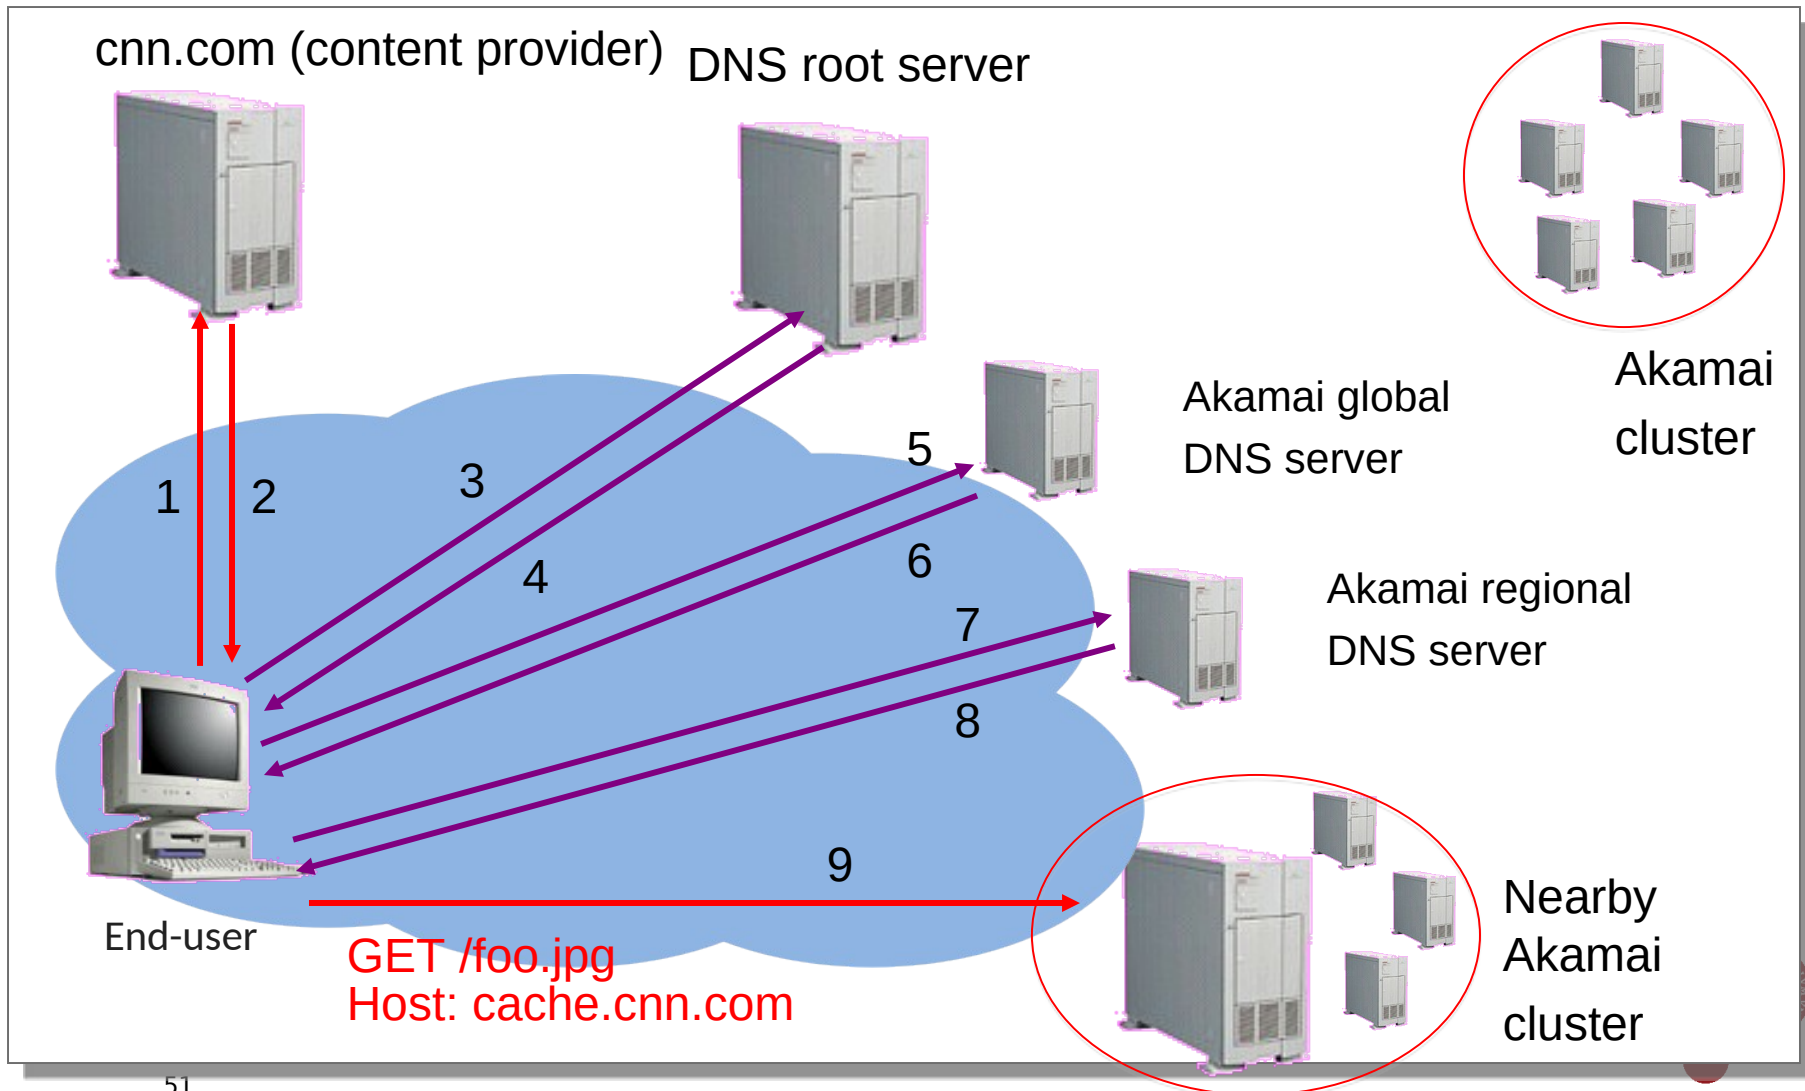
# How Akamai Works

Source: Freedman



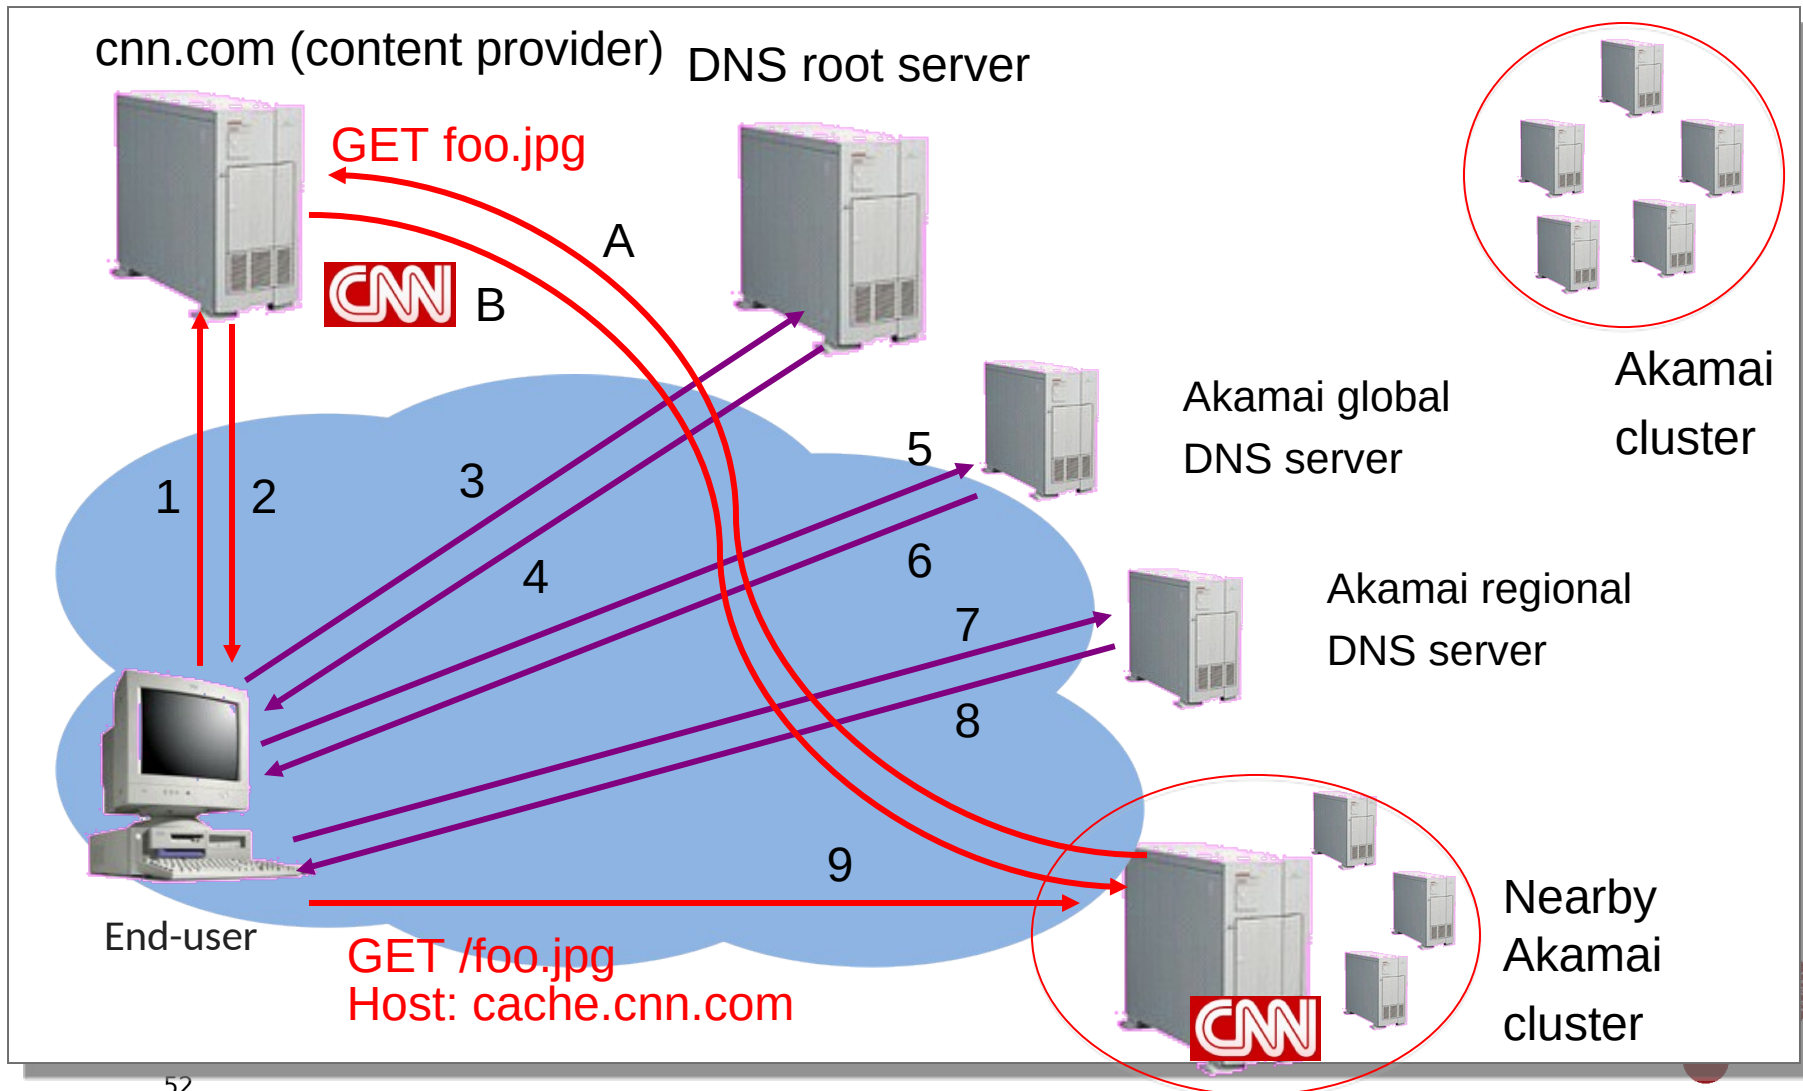
# How Akamai Works

Source: Freedman



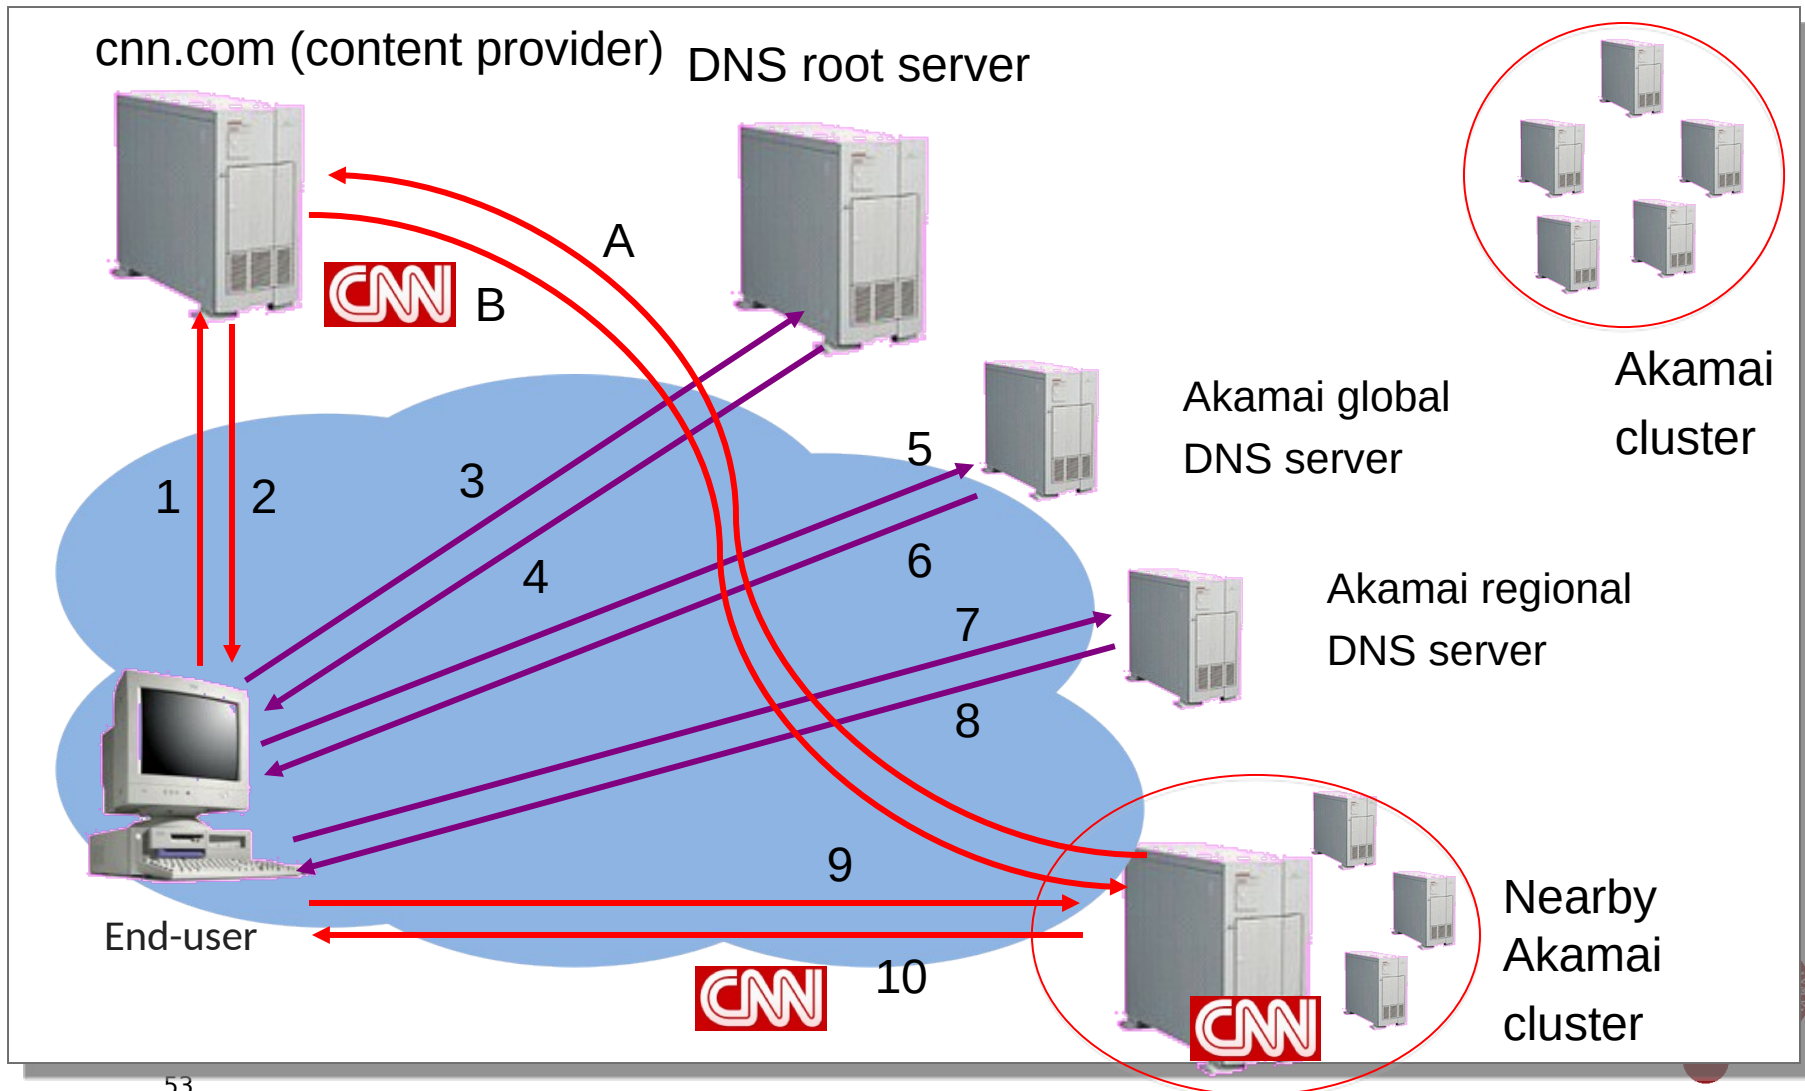
# How Akamai Works

Source: Freedman



# How Akamai Works

Source: Freedman

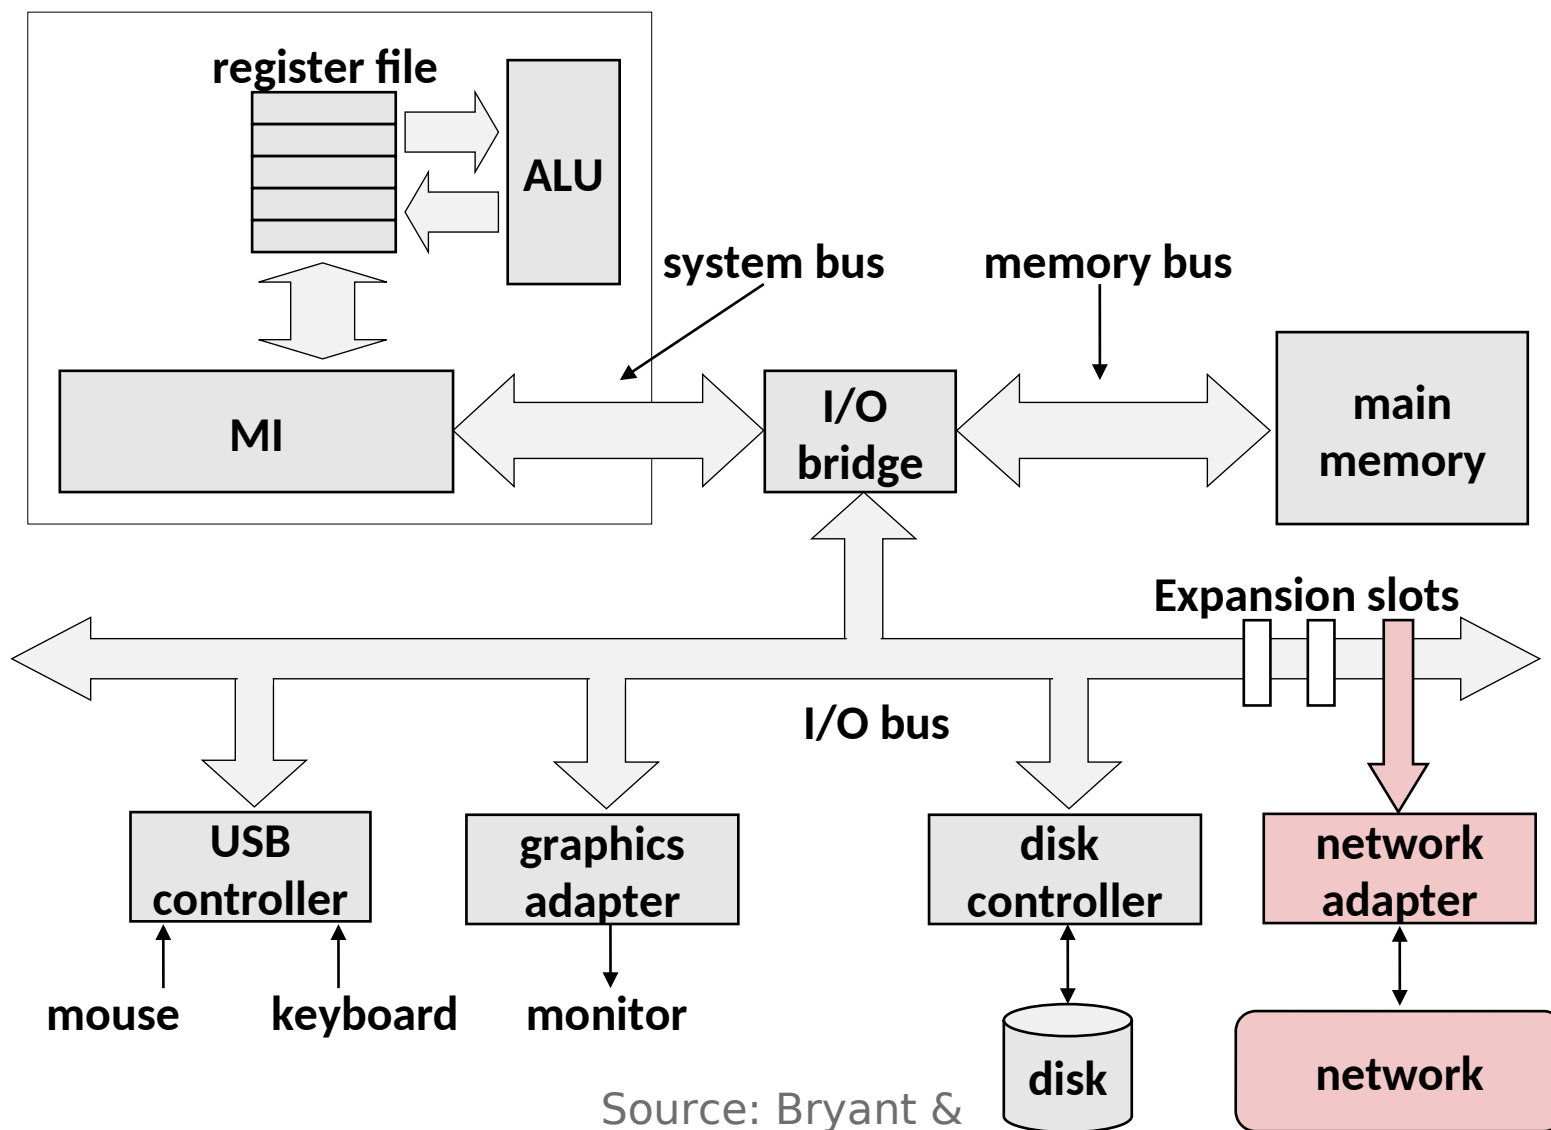


# Summary

- Network applications
  - Email, Web → more in textbook, Chapter 2!
- Socket abstraction
  - Communication between processes
  - Client / Server, Peer-to-Peer
- HTTP concepts
  - Web objects, request / response (pull)
  - Persistent connections, web proxies
  - Caching, content delivery



# What's next? More network programming in C ?



Source: Bryant & Halloran

